

Libraries

```
library(tidyverse)
# library(psych)           # describe()
library(DataExplorer)     # plot_missing() | drop_columns()
library(caret)            # nearZeroVar() | knnreg()
# library(inspectdf)      # inspect_cat() | show_plots()
# library(ggstance)       # geom_boxplot()
# library(corrplot)       # corrplot() | cor()
# library(ggpubr)         # ggscatter()
library(MASS)             # stepAIC()
library(regclass)         # vif()
# library(leaps)          # regsubsets()
library(ggplot2)          # ggplot()
library(glmtoolbox)       # hlttest()
# library(purrr)          # map()
library(GGally)           # ggcorr() | ggpairs()
library(lindia)           # gg_cooksd() | gg_scalelocation
library(gridExtra)        # grid.arrange
# library(FNN)            # knn.reg()
# library(Metrics)        # mse()
library(glmnet)           # cv.glmnet()
library(ROCR)             # prediction() | performance()
library(stats)            # logLik()
library(MLmetrics)        # LogLoss()
#####CLUSTERS#####
library(mvtnorm)
library(RColorBrewer)
library(pheatmap)
library(cluster)

library(jtools)           # interact_plot()
library(broom)            # augment()
```

Import Data

```
getwd()
```

```
## [1] "C:/Users/dnguy/Desktop/2 Applied Stats/Project 2/Statistics2-project-2"
```

```
df = read.csv("Bank_Personal_Loan_Modelling.csv")
```

EDA

```
str(df)
```

```
## 'data.frame':    5000 obs. of  14 variables:
##  $ ID              : int  1 2 3 4 5 6 7 8 9 10 ...
##  $ Age              : int  25 45 39 35 35 37 53 50 35 34 ...
##  $ Experience       : int  1 19 15 9 8 13 27 24 10 9 ...
##  $ Income           : int  49 34 11 100 45 29 72 22 81 180 ...
##  $ ZIP.Code         : int  91107 90089 94720 94112 91330 92121 91711 93943 90089 93023 ...
##  $ Family           : int  4 3 1 1 4 4 2 1 3 1 ...
##  $ CCAvg            : num  1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
##  $ Education        : int  1 1 1 2 2 2 2 3 2 3 ...
##  $ Mortgage         : int  0 0 0 0 0 155 0 0 104 0 ...
##  $ Personal.Loan    : int  0 0 0 0 0 0 0 0 0 1 ...
##  $ Securities.Account: int  1 1 0 0 0 0 0 0 0 0 ...
##  $ CD.Account       : int  0 0 0 0 0 0 0 0 0 0 ...
##  $ Online           : int  0 0 0 0 0 1 1 0 1 0 ...
##  $ CreditCard       : int  0 0 0 0 1 0 0 1 0 0 ...
```

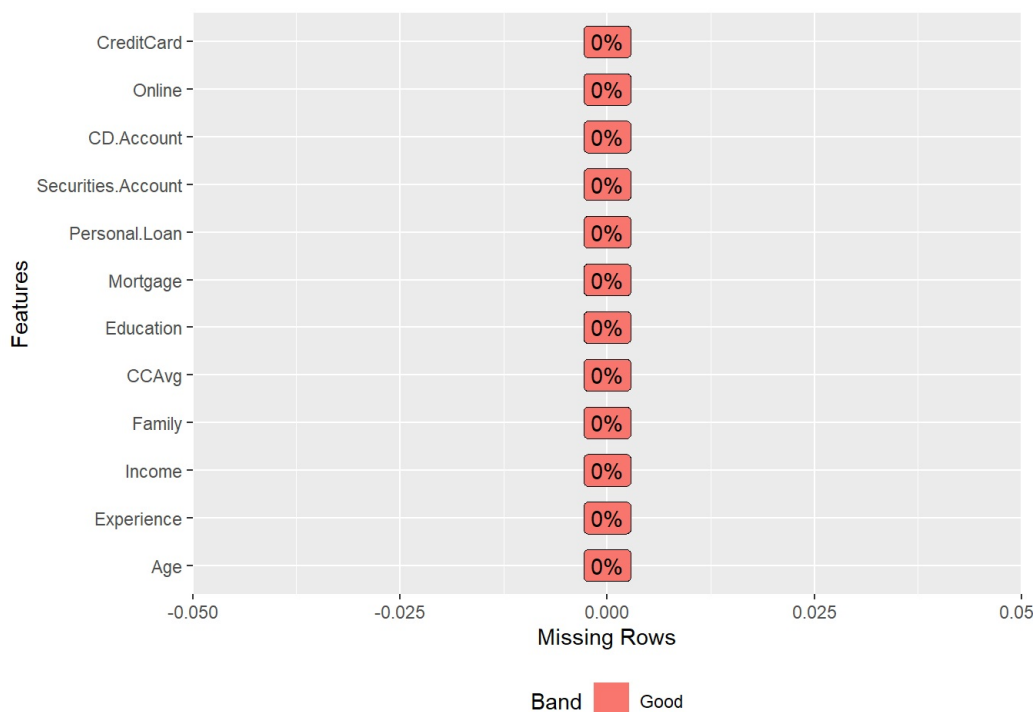
```
# Identification Columns (ID and ZIP.Code)
df = df[-c(1,5)]
str(df)
```

```
## 'data.frame': 5000 obs. of 12 variables:
## $ Age : int 25 45 39 35 35 37 53 50 35 34 ...
## $ Experience : int 1 19 15 9 8 13 27 24 10 9 ...
## $ Income : int 49 34 11 100 45 29 72 22 81 180 ...
## $ Family : int 4 3 1 1 4 4 2 1 3 1 ...
## $ CCAvg : num 1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
## $ Education : int 1 1 1 2 2 2 3 2 3 ...
## $ Mortgage : int 0 0 0 0 0 155 0 0 104 0 ...
## $ Personal.Loan : int 0 0 0 0 0 0 0 0 0 1 ...
## $ Securities.Account: int 1 1 0 0 0 0 0 0 0 0 ...
## $ CD.Account : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Online : int 0 0 0 0 0 1 1 0 1 0 ...
## $ CreditCard : int 0 0 0 0 1 0 0 1 0 0 ...
```

```
# Naturally Factor Variables
factor_vars = c("Family", "Education", "Personal.Loan",
               "Securities.Account", "CD.Account", "Online", "CreditCard")
df[factor_vars] = lapply(df[factor_vars], as.factor)
str(df)
```

```
## 'data.frame': 5000 obs. of 12 variables:
## $ Age : int 25 45 39 35 35 37 53 50 35 34 ...
## $ Experience : int 1 19 15 9 8 13 27 24 10 9 ...
## $ Income : int 49 34 11 100 45 29 72 22 81 180 ...
## $ Family : Factor w/ 4 levels "1","2","3","4": 4 3 1 1 4 4 2 1 3 1 ...
## $ CCAvg : num 1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
## $ Education : Factor w/ 3 levels "1","2","3": 1 1 1 2 2 2 2 3 2 3 ...
## $ Mortgage : int 0 0 0 0 0 155 0 0 104 0 ...
## $ Personal.Loan : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 2 ...
## $ Securities.Account: Factor w/ 2 levels "0","1": 2 2 1 1 1 1 1 1 1 1 ...
## $ CD.Account : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ Online : Factor w/ 2 levels "0","1": 1 1 1 1 1 2 2 1 2 1 ...
## $ CreditCard : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 2 1 1 ...
```

```
# missing values
plot_missing(df)
```



```
# near zero variance
nearZeroVar(df, names = TRUE)
```


```
## [1] "Mortgage"
```

```
#df = df[-c(nearZeroVar(df))] # Removed Mortgage
str(df)
```


```
## 'data.frame': 5000 obs. of 12 variables:
## $ Age : int 25 45 39 35 35 37 53 50 35 34 ...
## $ Experience : int 1 19 15 9 8 13 27 24 10 9 ...
## $ Income : int 49 34 11 100 45 29 72 22 81 180 ...
## $ Family : Factor w/ 4 levels "1","2","3","4": 4 3 1 1 4 2 1 3 1 ...
## $ CCAvg : num 1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
## $ Education : Factor w/ 3 levels "1","2","3": 1 1 1 2 2 2 2 3 2 3 ...
## $ Mortgage : int 0 0 0 0 0 155 0 0 104 0 ...
## $ Personal.Loan : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 2 ...
## $ Securities.Account : Factor w/ 2 levels "0","1": 2 2 1 1 1 1 1 1 1 ...
## $ CD.Account : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 ...
## $ Online : Factor w/ 2 levels "0","1": 1 1 1 1 1 2 2 1 2 1 ...
## $ CreditCard : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 2 1 1 ...
```

```
# multicollinearity
ggcorr(df, label = T)
```

```
## Warning in ggcorr(df, label = T): data in column(s) 'Family', 'Education',
## 'Personal.Loan', 'Securities.Account', 'CD.Account', 'Online', 'CreditCard' are
## not numeric and were ignored
```




```
## Age and Experience have correlation of 1
ggplot( df, aes(Experience, Age)) + geom_point()
```



```
ggplot(df, aes(Experience/Age)) + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
df = df %>% mutate(Experience2 = Experience/Age)
df = df[-c(1,2)] # getting rid of Age and Experience
str(df)
```

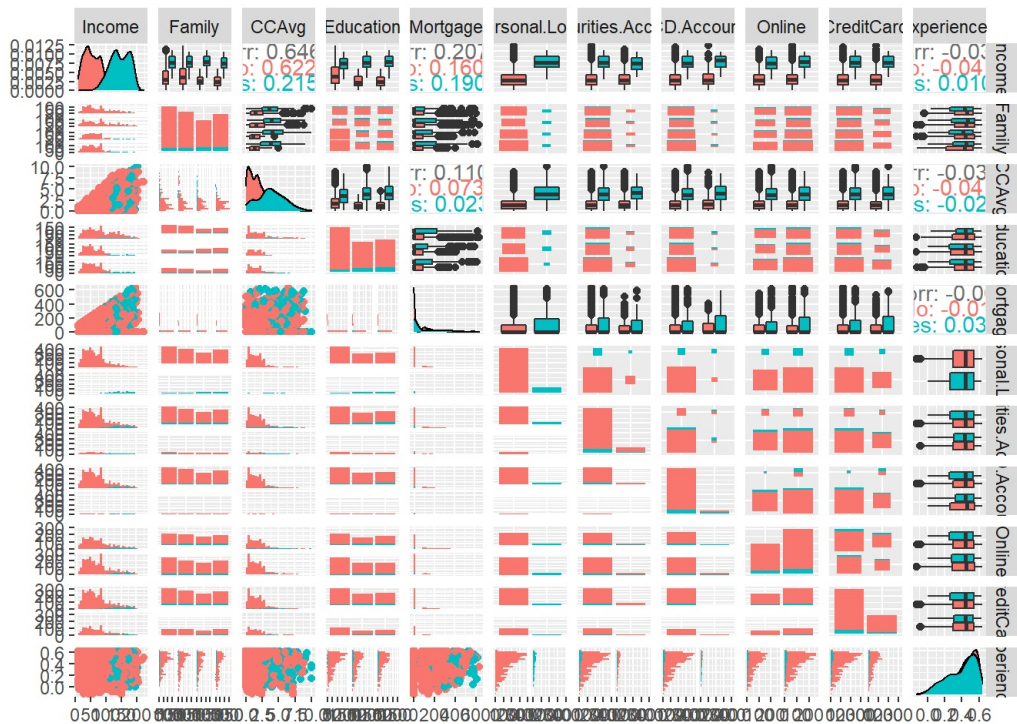
```
## 'data.frame':    5000 obs. of  11 variables:
## $ Income       : int  49 34 11 100 45 29 72 22 81 180 ...
## $ Family       : Factor w/ 4 levels "1","2","3","4": 4 3 1 1 4 2 1 3 1 ...
## $ CCAvg        : num  1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
## $ Education    : Factor w/ 3 levels "1","2","3": 1 1 1 2 2 2 2 3 2 3 ...
## $ Mortgage     : int   0 0 0 0 155 0 0 104 0 ...
## $ Personal.Loan : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 2 ...
## $ Securities.Account: Factor w/ 2 levels "0","1": 2 2 1 1 1 1 1 1 1 ...
## $ CD.Account   : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 ...
## $ Online       : Factor w/ 2 levels "0","1": 1 1 1 1 1 2 2 1 2 1 ...
## $ CreditCard   : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 2 1 1 ...
## $ Experience2   : num  0.04 0.422 0.385 0.257 0.229 ...
```

```
ggcorr(df, label = T)
```

```
## Warning in ggcorr(df, label = T): data in column(s) 'Family', 'Education',  
## 'Personal.Loan', 'Securities.Account', 'CD.Account', 'Online', 'CreditCard' are  
## not numeric and were ignored
```




```
# pairs plots  
#newAuto$mpg<- factor(ifelse(Auto$mpg>median(Auto$mpg),"High","Low"),levels=c("Low","High")) # used for numeric  
outcome into categorical outcome (using median)  
# kept for future reference  
levels(df$Personal.Loan) = c("No", "Yes")  
ggpairs(df, aes(colour = Personal.Loan))
```



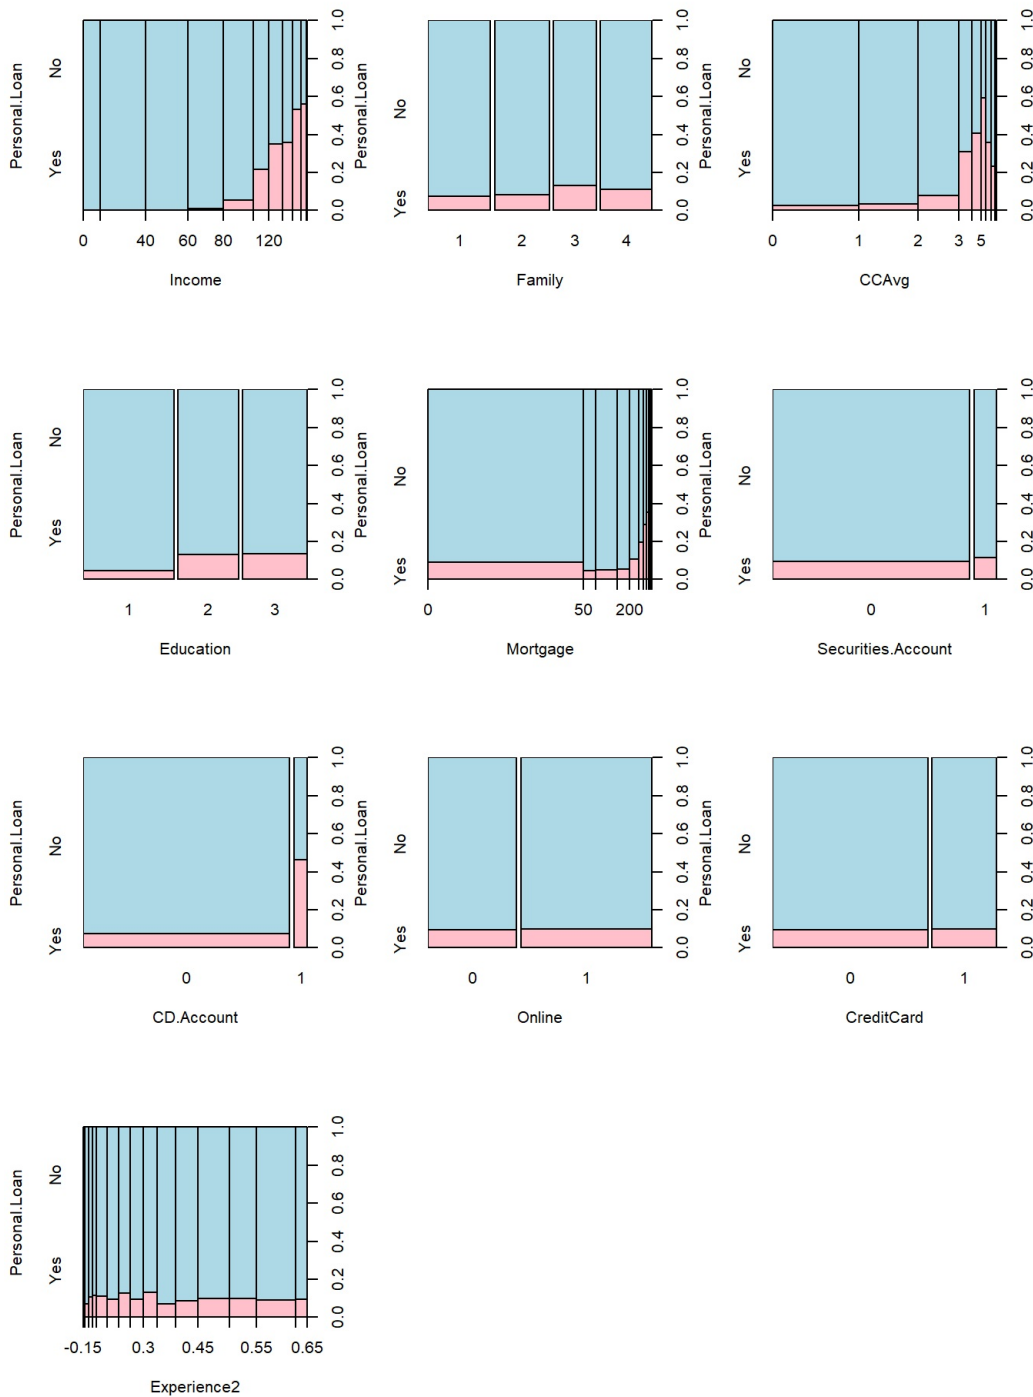
```
## Using the trick of already knowing what my stepwise logistic regression model consists of in terms of coefficients (which are Income, Family, CCAvg, Education, Securities.Account, CD.Account, Online, and CreditCard) I can pretend to say that the following variables can be considered in our model for Objective 1 to predict whether if a customer will accept a personal loan offer or not.
```

```
## We can see that, for variables with multiple levels, the levels with even a slight change compared to the reference level (1st level) are found as significant to our stepwise logistic regression model.
```

```
# This determines green is yes.  
plot(df$Personal.Loan, col= c("lightblue","pink"))
```



```
par(mfrow=c(2,3))  
plot(Personal.Loan ~ ., data = df, col= c("pink","lightblue"))
```



EDA: Exploring Interactions

```
#interact_plot()
```

EDA: Heatmaps (Unit 13)

EDA: Cluster Analysis (Unit 13)

Train Test Split

```
set.seed(123)

split = sample(nrow(df), nrow(df)*0.7)

train = df[split,]
test = df[-split,]
```

Objective 1: Logistic Regression Model

```
premodel = glm(Personal.Loan ~ ., data = train, family = "binomial")
```

```
# feature selection - stepwise  
stepAIC(premodel, direction = "both")
```

```
## Start: AIC=820.41  
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +  
## Securities.Account + CD.Account + Online + CreditCard + Experience2  
##  
##           Df Deviance    AIC  
## - Experience2      1   792.42  818.42  
## - Mortgage        1   794.06  820.06  
## <none>              792.41  820.41  
## - Securities.Account 1   797.09  823.09  
## - CCAvg            1   798.20  824.20  
## - CreditCard       1   805.06  831.06  
## - Online           1   809.23  835.23  
## - Family           3   865.63  887.63  
## - CD.Account       1   879.79  905.79  
## - Education        2  1083.03 1107.03  
## - Income           1  1358.72 1384.72  
##  
## Step: AIC=818.42  
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +  
## Securities.Account + CD.Account + Online + CreditCard  
##  
##           Df Deviance    AIC  
## - Mortgage        1   794.07  818.07  
## <none>              792.42  818.42  
## + Experience2      1   792.41  820.41  
## - Securities.Account 1   797.10  821.10  
## - CCAvg            1   798.32  822.32  
## - CreditCard       1   805.07  829.07  
## - Online           1   809.25  833.25  
## - Family           3   865.64  885.64  
## - CD.Account       1   879.96  903.96  
## - Education        2  1083.11 1105.11  
## - Income           1  1358.75 1382.75  
##  
## Step: AIC=818.07  
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +  
## CD.Account + Online + CreditCard  
##  
##           Df Deviance    AIC  
## <none>              794.07  818.07  
## + Mortgage        1   792.42  818.42  
## + Experience2      1   794.06  820.06  
## - Securities.Account 1   798.62  820.62  
## - CCAvg            1   799.43  821.43  
## - CreditCard       1   807.04  829.04  
## - Online           1   810.78  832.78  
## - Family           3   867.01  885.01  
## - CD.Account       1   882.58  904.58  
## - Education        2  1083.88 1103.88  
## - Income           1  1387.09 1409.09
```

```
##  
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +  
## Securities.Account + CD.Account + Online + CreditCard, family = "binomial",  
## data = train)  
##  
## Coefficients:  
## (Intercept)      Income      Family2  
## -12.54761      0.06381     -0.20616  
## Family3      Family4      CCAvg  
## 1.92186      1.39674      0.12679  
## Education2      Education3 Securities.Account1  
## 4.01753      4.17725     -0.71361  
## CD.Account1      Online1      CreditCard1  
## 3.54524     -0.81711     -0.88671  
##  
## Degrees of Freedom: 3499 Total (i.e. Null); 3488 Residual  
## Null Deviance: 2164  
## Residual Deviance: 794.1 AIC: 818.1
```



```
modell1 = glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +  
  Securities.Account + CD.Account + Online + CreditCard, family = "binomial",  
  data = train)
```

Hypothesis Testing

```
summary(modell1)
```

```
##  
## Call:  
## glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +  
##   Securities.Account + CD.Account + Online + CreditCard, family = "binomial",  
##   data = train)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max  
## -2.9123  -0.1813  -0.0649  -0.0194   4.1340  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)    -12.547610    0.663989  -18.897 < 2e-16 ***  
## Income           0.063807    0.003717   17.167 < 2e-16 ***  
## Family2        -0.206161    0.281441   -0.733 0.463853  
## Family3         1.921856    0.295621    6.501 7.97e-11 ***  
## Family4         1.396735    0.290604    4.806 1.54e-06 ***  
## CCAvg           0.126789    0.055110    2.301 0.021412 *  
## Education2      4.017534    0.334135   12.024 < 2e-16 ***  
## Education3      4.177245    0.333415   12.529 < 2e-16 ***  
## Securities.Account1 -0.713607    0.348807   -2.046 0.040771 *  
## CD.Account1     3.545241    0.403945    8.777 < 2e-16 ***  
## Online1        -0.817114    0.201944   -4.046 5.20e-05 ***  
## CreditCard1    -0.886710    0.258336   -3.432 0.000598 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##    Null deviance: 2163.69  on 3499  degrees of freedom  
## Residual deviance: 794.07  on 3488  degrees of freedom  
## AIC: 818.07  
##  
## Number of Fisher Scoring iterations: 8
```

As the p-values of all variables used in modell1, aside from Family2, are all less than 0.05, none of them are insignificant in our logistic regression model.

Criterion

```
AIC(modell1)      # AIC = 818.07
```

```
## [1] 818.0738
```

```
BIC(modell1)      # BIC = 892
```

```
## [1] 892
```

Verify Predictions Manually

```
# Holding the upcoming predictions accountable  
prop.table(table(df$Personal.Loan))
```

```
##  
##      No    Yes  
## 0.904 0.096
```

```
prop.table(table(train$Personal.Loan))
```

```
##
##           No           Yes
## 0.90714286 0.09285714
```

```
prop.table(table(test$Personal.Loan))
```

```
##
##           No           Yes
## 0.8966667 0.1033333
```

```
# This means that,
# it is preferred that our predictions are 90% no loan and 10% yes loan.
```

The general idea is, for a bank problem like this where we are trying to find profit from the highest number of customers who will accept a personal loan offer as we can, we want to have an as-low-as-possible chance of predicting customers saying no but they actually do want to say yes because, not calling an interested customer will cost us valuable profits. However, calling a disinterested customer will not hurt that much where they will simply assume that it's a cold call. Unless our decisions mean that the bank can forcibly and automatically give a customer a loan despite them not being interested in a loan, or rather a more realistic example like using software to determine a patient to have cancer even though they do not, and that patient will wastefully go through a surgery process, like an actionable decision from our predictions, we should be fine with a low specificity (with the other proportion of low specificity meaning a high chance of predicting yes to people saying no, which is perfectly OK and can be overlooked).

```
pred.step = predict(model1, test, type = "response")

step.cutoff = 0.426
class.step = as.factor(if_else(pred.step < step.cutoff, "No", "Yes"))
#pred = as.factor(if_else(pred < 0.3, 0, 1))
prop.table(table(class.step))
```

```
## class.step
##           No           Yes
## 0.914 0.086
```

```
# Confusion Matrix
confusionMatrix(class.step, test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1334  37
##           Yes  11 118
##
##           Accuracy : 0.968
##           95% CI : (0.9578, 0.9763)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8135
##
##           Mcnemar's Test P-Value : 0.000308
##
##           Sensitivity : 0.9918
##           Specificity : 0.7613
##           Pos Pred Value : 0.9730
##           Neg Pred Value : 0.9147
##           Prevalence : 0.8967
##           Detection Rate : 0.8893
##           Detection Prevalence : 0.9140
##           Balanced Accuracy : 0.8766
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.426
# Accuracy = 0.968
# Sensitivity = 0.7613
# Specificity = 0.9919
```

Assumptions

```
# Linearity
## Predict the probability (p) of personal loan offer
probabilities <- predict(modell, type = "response")
length(probabilities)
```

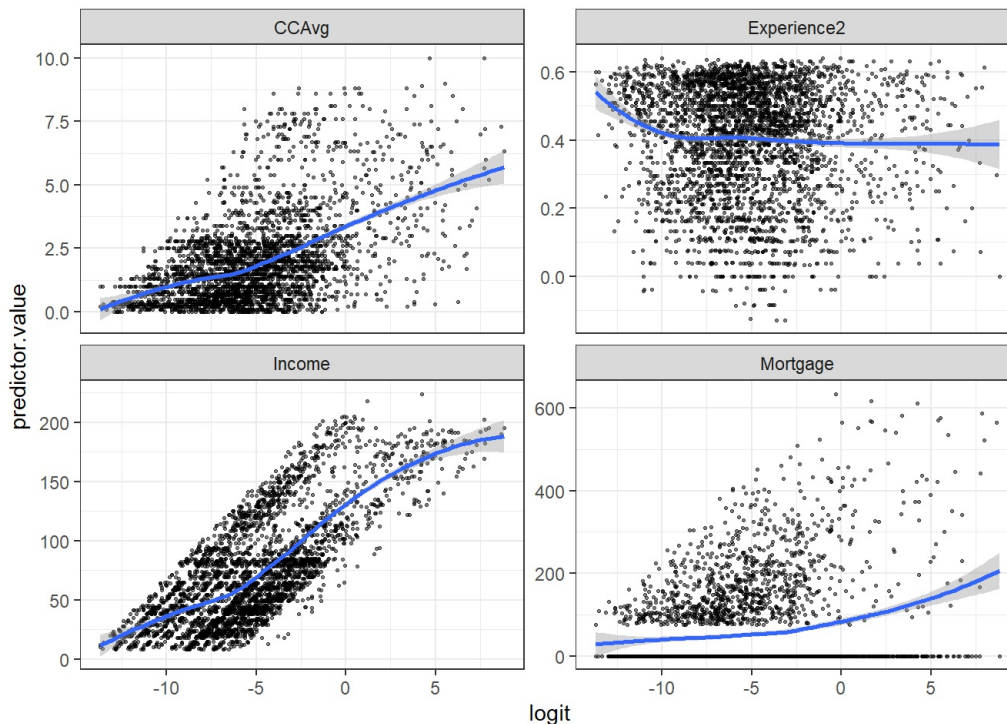
```
## [1] 3500
```

```
step.cutoff = 0.3
predicted.classes <- ifelse(probabilities > step.cutoff, "Yes", "No")
head(predicted.classes)
```



```
## 2463 2511 2227 526 4291 2986
## "No" "No" "No" "No" "No" "No"
```

```
## Select only numeric predictors
mydata <- train %>% select_if(is.numeric)
predictors <- colnames(mydata)
## Bind the logit and tidying the data for plot
mydata <- mydata %>%
  mutate(logit = log(probabilities/(1-probabilities))) %>%
  gather(key = "predictors", value = "predictor.value", -logit)
## Create scatter plots
ggplot(mydata, aes(logit, predictor.value))+
  geom_point(size = 0.5, alpha = 0.5) +
  geom_smooth(method = "loess") +
  theme_bw() +
  facet_wrap(~predictors, scales = "free_y")
```

```
## `geom_smooth()` using formula 'y ~ x'
```




```
# Influential Points
par(mfrow = c(1, 2))
## Cook's Distance Plot
plot(modell, 4, 3)
## Standardized Residuals vs Leverage
plot(modell, 5, 3)
```



```
par(mfrow = c(1, 1))
## Extract model results
model.data <- augment(model1) %>%
  mutate(index = 1:n())
model.data %>% top_n(3, .cooksd)
```

```
## # A tibble: 3 x 17
##   .rownames Persona~1 Income Family CCAvg Educa~2 Secur~3 CD.Ac~4 Online Credi~5
##   <chr>      <fct>      <int> <fct>  <dbl> <fct>  <fct>  <fct>  <fct>  <fct>
## 1 2540      Yes         98 1     4.2 1     1     1     0     0
## 2 350       Yes         60 2     3   1     0     0     0     0
## 3 2346      Yes         89 1     4.1 1     0     1     1     0
## # ... with 7 more variables: .fitted <dbl>, .resid <dbl>, .std.resid <dbl>,
## #   .hat <dbl>, .sigma <dbl>, .cooksd <dbl>, index <int>, and abbreviated
## #   variable names 1: Personal.Loan, 2: Education, 3: Securities.Account,
## #   4: CD.Account, 5: CreditCard
## # i Use `colnames()` to see all variable names
```

```
ggplot(model.data, aes(index, .std.resid)) +
  geom_point(aes(color = Personal.Loan), alpha = .5) +
  theme_bw()
```



```
## Culprit Outlier Observations
outliers = model.data %>% filter(abs(.std.resid) > 3)
outliers$.rownames
```

```
## [1] "1127" "1070" "976" "350" "2159"
```

```
# Multicollinearity
VIF(model1)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## Income          2.940809  1      1.714879
## Family          1.529409  3      1.073381
## CCAvg           1.516750  1      1.231564
## Education       2.323075  2      1.234570
## Securities.Account 1.291648  1      1.136507
## CD.Account      1.936714  1      1.391659
## Online          1.143566  1      1.069376
## CreditCard      1.383602  1      1.176266
```

Interpretations and Confidence Intervals

```
# Coefficients
coef(model1)
```

```
##      (Intercept)          Income          Family2          Family3
##      -12.54761027      0.06380744      -0.20616070      1.92185637
##           Family4          CCAvg          Education2          Education3
##           1.39673534      0.12678909      4.01753412      4.17724518
## Securities.Account1      CD.Account1          Online1          CreditCard1
##           -0.71360741      3.54524089      -0.81711404      -0.88670963
```

```
# interpret as log odds & confidence intervals
format(exp(cbind("Odds Ratio" = coef(model1),
                 confint.default(model1, level = 0.95))),
       scientific = F)
```

```
##      Odds Ratio      2.5 %      97.5 %
## (Intercept) " 0.0000035533836" " 0.0000009670612" " 0.0000130566048"
## Income      " 1.0658871357374" " 1.0581505454691" " 1.0736802915192"
## Family2     " 0.8137023043692" " 0.4687078501608" " 1.4126314289566"
## Family3     " 6.8336324635477" " 3.8284108534361" " 12.1978895250861"
## Family4     " 4.0419826931018" " 2.2868175478207" " 7.1442621677030"
## CCAvg       " 1.1351775688223" " 1.0189519542112" " 1.2646603281260"
## Education2  " 55.5639227354387" " 28.8653266780218" "106.9570264763547"
## Education3  " 65.1860299588602" " 33.9118239040492" "125.3019747277599"
## Securities.Account1 " 0.4898738322890" " 0.2472743352565" " 0.9704863681571"
## CD.Account1 " 34.6480306876299" " 15.6978347366024" " 76.4746253654833"
## Online1     " 0.4417045599866" " 0.2973287142396" " 0.6561859281298"
## CreditCard1 " 0.4120091864564" " 0.2483194731810" " 0.6836015216606"
```

```
# Holding all other variables constant,
### an increase of $1,000 in a customer's income is associated with an increase of 6.58871% in the odds of them accepting a personal loan offer.
### customers with a family size of 2 have around 0.814 times the odds of accepting a personal loan offer than those who don't.
### ...
### customers with a securities account have a .49 times the odds of those who don't of accepting a personal loan offer.
### ...
```

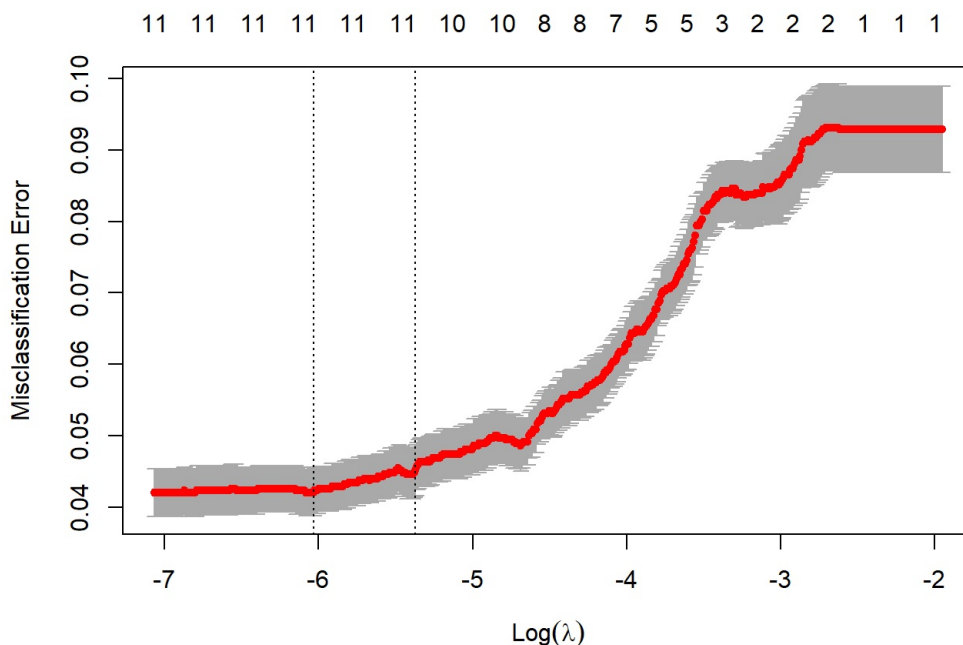
Objective 1: LASSO Penalized Logistic Regression Model

```
str(train)
```

```
## 'data.frame': 3500 obs. of 11 variables:
## $ Income : int 23 52 98 79 95 63 91 143 59 38 ...
## $ Family : Factor w/ 4 levels "1","2","3","4": 3 4 1 2 2 4 1 3 3 1 ...
## $ CCAvg : num 0.4 1.3 5.4 2.8 0 3.6 0.1 2.9 0.9 1.5 ...
## $ Education : Factor w/ 3 levels "1","2","3": 1 2 1 1 3 3 2 3 3 2 ...
## $ Mortgage : int 0 0 0 179 0 0 199 0 199 116 ...
## $ Personal.Loan : Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 1 2 1 ...
## $ Securities.Account: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 2 ...
## $ CD.Account : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ Online : Factor w/ 2 levels "0","1": 2 2 2 1 2 1 2 2 2 1 ...
## $ CreditCard : Factor w/ 2 levels "0","1": 1 2 1 1 1 1 2 1 1 2 ...
## $ Experience2 : num 0.538 0.613 0.04 0.594 0.636 ...
```

```
dat.train.x = model.matrix(Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account + CD.Account + Online + CreditCard + Experience2, train)
dat.train.y = train$Personal.Loan

cvfit = cv.glmnet(dat.train.x, dat.train.y, family = "binomial", type.measure = "class", nlambda = 1000)
plot(cvfit)
```



```
coef(cvfit, s = "lambda.min")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -10.67329812
## (Intercept) .
## Income      0.05403179
## Family2     -0.14483106
## Family3      1.52190649
## Family4      1.05043583
## CCAvg        0.08593637
## Education2   3.15116632
## Education3   3.26676559
## Securities.Account1 -0.29480894
## CD.Account1  2.68760355
## Online1      -0.51980937
## CreditCard1  -0.52133702
## Experience2  .
```

```
# CV misclassification error rate is little below .10
cvfit$cvm[which(cvfit$lambda==cvfit$lambda.min)]
```

```
## [1] 0.042
```

```
# Optimal penalty
cvfit$lambda.min
```

```
## [1] 0.002395327
```

```
# For final model predictions go ahead and refit lasso using entire data set
LASSOmodel = glmnet(dat.train.x, dat.train.y, family = "binomial", lambda=cvfit$lambda.min)
coef(LASSOmodel, s = "lambda.min")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                -10.67454150
## (Intercept)                  .
## Income                      0.05403933
## Family2                     -0.14480803
## Family3                      1.52200412
## Family4                      1.05046564
## CCAvg                       0.08592698
## Education2                   3.15171466
## Education3                   3.26728389
## Securities.Account1         -0.29482677
## CD.Account1                  2.68779906
## Online1                     -0.51987111
## CreditCard1                 -0.52136272
## Experience2                  .
```

```
# Predict
dat.test.x = model.matrix(Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account + CD.Account +
Online + CreditCard + Experience2, test)
fit.pred.lasso = predict(LASSOmodel, newx = dat.test.x, type = "response")

LASSO.cutoff = 0.44
class.lasso = as.factor(if_else(fit.pred.lasso < LASSO.cutoff, "No", "Yes"))

# Confusion Matrix for Lasso
conf.lasso = table(class.lasso, test$Personal.Loan)
conf.lasso
```

```
##
## class.lasso   No   Yes
##              No 1340   50
##              Yes   5  105
```

```
# Accuracy of LASSO
sum(diag(conf.lasso))/sum(conf.lasso)
```

```
## [1] 0.9633333
```

```
# Sensitivity & Specificity of LASSO
cm = confusionMatrix(class.lasso, test$Personal.Loan)
cm$byClass
```

```
##          Sensitivity      Specificity      Pos Pred Value
##          0.9962825        0.6774194        0.9640288
##          Neg Pred Value      Precision          Recall
##          0.9545455        0.9640288        0.9962825
##          F1          Prevalence      Detection Rate
##          0.9798903        0.8966667        0.8933333
## Detection Prevalence      Balanced Accuracy
##          0.9266667        0.8368509
```

```
# Threshold = 0.44
# Accuracy = 0.9633
# Sensitivity = 0.6774
# Specificity = 0.9963
```

Objective 1: Erin's Model based on Intuition

```
mod.erin = glm(formula = Personal.Loan ~ Income + Family + Education + CD.Account + CreditCard,
               family = "binomial", data = train)
summary(mod.erin)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ Income + Family + Education + CD.Account +
##      CreditCard, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -2.7993  -0.1856  -0.0677  -0.0215   4.2636
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -12.902065   0.650445 -19.836 < 2e-16 ***
## Income       0.066222   0.003506  18.888 < 2e-16 ***
## Family2     -0.160351   0.277510  -0.578  0.56339
## Family3      1.986369   0.289837   6.853 7.21e-12 ***
## Family4      1.405825   0.283468   4.959 7.07e-07 ***
## Education2   3.875382   0.324913  11.927 < 2e-16 ***
## Education3   4.042373   0.322560  12.532 < 2e-16 ***
## CD.Account1  2.829847   0.333789   8.478 < 2e-16 ***
## CreditCard1 -0.703303   0.245175  -2.869  0.00412 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2163.69  on 3499  degrees of freedom
## Residual deviance:  818.88  on 3491  degrees of freedom
## AIC: 836.88
##
## Number of Fisher Scoring iterations: 8
```

```
# Criterion
AIC(mod.erin)          # AIC = 836.8792
```

```
## [1] 836.8792
```

```
BIC(mod.erin)          # BIC = 892.3239
```

```
## [1] 892.3239
```

```
pred.erin = predict(mod.erin, test, type = "response")
erin.cutoff = 0.5
class.erin = as.factor(if_else(pred.erin < erin.cutoff, "No", "Yes"))
prop.table(table(class.erin))
```

```
## class.erin
##      No   Yes
## 0.928 0.072
```

```
confusionMatrix(class.erin, test$Personal.Loan)
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1338  54
##           Yes   7 101
##
##           Accuracy : 0.9593
##           95% CI : (0.9481, 0.9688)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7466
##
##           McNemar's Test P-Value : 3.869e-09
##
##           Sensitivity : 0.9948
##           Specificity : 0.6516
##           Pos Pred Value : 0.9612
##           Neg Pred Value : 0.9352
##           Prevalence : 0.8967
##           Detection Rate : 0.8920
##           Detection Prevalence : 0.9280
##           Balanced Accuracy : 0.8232
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.5
# Accuracy = 0.9593
# Sensitivity = 0.6516
# Specificity = 0.9948
```

Objective 1: Origin Model (Income Only)

```
model_income = glm(formula = Personal.Loan ~ Income, family = "binomial", data = train)
summary(model_income)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ Income, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0919  -0.3066  -0.1796  -0.1166   2.7881
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.03738    0.21904  -27.56  <2e-16 ***
## Income       0.03619    0.00163   22.20  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2163.7  on 3499  degrees of freedom
## Residual deviance: 1407.4  on 3498  degrees of freedom
## AIC: 1411.4
##
## Number of Fisher Scoring iterations: 6
```

```
pred.income = predict(model_income, test, type = "response")

income.cutoff = 0.3
class.income = as.factor(if_else(pred.income < income.cutoff, "No", "Yes"))
prop.table(table(class.income))
```

```
## class.income
##           No           Yes
## 0.8873333 0.1126667
```

```
confusionMatrix(class.income, test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1258  73
##           Yes  87  82
##
##           Accuracy : 0.8933
##           95% CI : (0.8766, 0.9085)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : 0.6827
##
##           Kappa : 0.4465
##
## Mcnemar's Test P-Value : 0.3041
##
##           Sensitivity : 0.9353
##           Specificity : 0.5290
##           Pos Pred Value : 0.9452
##           Neg Pred Value : 0.4852
##           Prevalence : 0.8967
##           Detection Rate : 0.8387
##           Detection Prevalence : 0.8873
##           Balanced Accuracy : 0.7322
##
##           'Positive' Class : No
##
```

```
           # Threshold = 0.3
           # Accuracy = 0.8933
           # Sensitivity = 0.5290
           # Specificity = 0.9353
# Criterion
AIC(model_income)           # AIC = 1411.45
```

```
## [1] 1411.45
```

```
BIC(model_income)           # BIC = 1423.771
```

```
## [1] 1423.771
```

Comparing ROCR Curves


```
# Stepwise
pred_prob = predict(model1, test, type = "response")
test_label = df[-split, "Personal.Loan"]
results.step = prediction(pred_prob, test_label)
roc.step = performance(results.step, measure = "tpr", x.measure = "fpr")

# LASSO
results.lasso = prediction(fit.pred.lasso,
                           test$Personal.Loan,
                           label.ordering=c("No", "Yes"))
roc.lasso = performance(results.lasso, measure = "tpr", x.measure = "fpr")

# Erin's Intuition
results.erin = prediction(pred.erin, test_label)
roc.erin = performance(results.erin, measure = "tpr", x.measure = "fpr")

# Origin (Income Only)
results.income = prediction(pred.income, test_label)
roc.income = performance(results.income, measure = "tpr", x.measure = "fpr")

plot(roc.step, col = "red", xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.lasso, col = "green", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.erin, col = "blue", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.income, col = "pink", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
legend("bottomright", legend = c("Stepwise", "Lasso", "Erin", "Origin (Income Only)"),
      col = c("red", "green", "blue", "pink"),
      lty=1, lwd=1)
```



```
#abline(a=0, b= 1)
#abline(a=1, b= -1)

# Stepwise seems to be the better performing model according to the above ROC curves.
```

Objective 2: Adding Complexity

```
model.poly.income2 = glm(formula = Personal.Loan ~ poly(Income, 2) + Family + CCAvg + Education + Securities.Acco
unt + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.income2 = predict(model.poly.income2, test, type = "response")
poly.income2.cutoff = 0.55
class.poly.income2 = as.factor(if_else(pred.poly.income2 < poly.income2.cutoff, "No", "Yes"))
prop.table(table(class.poly.income2))
```

```
## class.poly.income2
##      No      Yes
## 0.9113333 0.0886667
```

```
confusionMatrix(class.poly.income2, test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1334  33
##           Yes  11 122
##
##           Accuracy : 0.9707
##           95% CI : (0.9608, 0.9786)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8311
##
##           McNemar's Test P-Value : 0.001546
##
##           Sensitivity : 0.9918
##           Specificity : 0.7871
##           Pos Pred Value : 0.9759
##           Neg Pred Value : 0.9173
##           Prevalence : 0.8967
##           Detection Rate : 0.8893
##           Detection Prevalence : 0.9113
##           Balanced Accuracy : 0.8895
##
##           'Positive' Class : No
##
```

```
results.poly.income2 = prediction(pred.poly.income2, test_label)
roc.poly.income2 = performance(results.poly.income2, measure = "tpr", x.measure = "fpr")

model.poly.income3 = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + CCAvg + Education + Securities.Account + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.income3 = predict(model.poly.income3, test, type = "response")
poly.income3.cutoff = 0.55
class.poly.income3 = as.factor(if_else(pred.poly.income3 < poly.income3.cutoff, "No", "Yes"))
prop.table(table(class.poly.income3))
```

```
## class.poly.income3
##           No           Yes
## 0.91066667 0.08933333
```

```
confusionMatrix(class.poly.income3, test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1333  33
##           Yes  12 122
##
##           Accuracy : 0.97
##           95% CI : (0.9601, 0.978)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8278
##
##           McNemar's Test P-Value : 0.002869
##
##           Sensitivity : 0.9911
##           Specificity : 0.7871
##           Pos Pred Value : 0.9758
##           Neg Pred Value : 0.9104
##           Prevalence : 0.8967
##           Detection Rate : 0.8887
##           Detection Prevalence : 0.9107
##           Balanced Accuracy : 0.8891
##
##           'Positive' Class : No
##
```

```

results.poly.income3 = prediction(pred.poly.income3, test_label)
roc.poly.income3 = performance(results.poly.income3, measure = "tpr", x.measure = "fpr")

model.poly.CCAvg2 = glm(formula = Personal.Loan ~ Income + Family + poly(CCAvg, 2) + Education + Securities.Account + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.CCAvg2 = predict(model.poly.CCAvg2, test, type = "response")
poly.CCAvg2.cutoff = 0.55
class.poly.CCAvg2 = as.factor(if_else(pred.poly.CCAvg2 < poly.CCAvg2.cutoff, "No", "Yes"))
prop.table(table(class.poly.CCAvg2))

```

```

## class.poly.CCAvg2
##           No           Yes
## 0.92133333 0.07866667

```

```

confusionMatrix(class.poly.CCAvg2, test$Personal.Loan)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1339  43
##           Yes   6 112
##
##           Accuracy : 0.9673
##           95% CI : (0.957, 0.9757)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8029
##
##           Mcnemar's Test P-Value : 2.706e-07
##
##           Sensitivity : 0.9955
##           Specificity : 0.7226
##           Pos Pred Value : 0.9689
##           Neg Pred Value : 0.9492
##           Prevalence : 0.8967
##           Detection Rate : 0.8927
##           Detection Prevalence : 0.9213
##           Balanced Accuracy : 0.8591
##
##           'Positive' Class : No
##

```

```

results.poly.CCAvg2 = prediction(pred.poly.CCAvg2, test_label)
roc.poly.CCAvg2 = performance(results.poly.CCAvg2, measure = "tpr", x.measure = "fpr")

model.poly.CCAvg3 = glm(formula = Personal.Loan ~ Income + Family + poly(CCAvg, 3) + Education + Securities.Account + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.CCAvg3 = predict(model.poly.CCAvg3, test, type = "response")
poly.CCAvg3.cutoff = 0.55
class.poly.CCAvg3 = as.factor(if_else(pred.poly.CCAvg3 < poly.CCAvg3.cutoff, "No", "Yes"))
prop.table(table(class.poly.CCAvg3))

```

```

## class.poly.CCAvg3
##           No           Yes
## 0.91933333 0.08066667

```

```

confusionMatrix(class.poly.CCAvg3, test$Personal.Loan)

```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1337  42
##           Yes   8 113
##
##           Accuracy : 0.9667
##           95% CI : (0.9563, 0.9752)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8008
##
##           McNemar's Test P-Value : 3.058e-06
##
##           Sensitivity : 0.9941
##           Specificity : 0.7290
##           Pos Pred Value : 0.9695
##           Neg Pred Value : 0.9339
##           Prevalence : 0.8967
##           Detection Rate : 0.8913
##           Detection Prevalence : 0.9193
##           Balanced Accuracy : 0.8615
##
##           'Positive' Class : No
##
```

```
results.poly.CCAvg3 = prediction(pred.poly.CCAvg3, test_label)
roc.poly.CCAvg3 = performance(results.poly.CCAvg3, measure = "tpr", x.measure = "fpr")

model.poly.both2 = glm(formula = Personal.Loan ~ poly(Income, 2) + Family + poly(CCAvg, 2) + Education + Securities.Account + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.both2 = predict(model.poly.both2, test, type = "response")
poly.both2.cutoff = 0.55
class.poly.both2 = as.factor(if_else(pred.poly.both2 < poly.both2.cutoff, "No", "Yes"))
prop.table(table(class.poly.both2))
```

```
## class.poly.both2
##   No  Yes
## 0.91 0.09
```

```
confusionMatrix(class.poly.both2, test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1334  31
##           Yes  11 124
##
##           Accuracy : 0.972
##           95% CI : (0.9623, 0.9797)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.8398
##
##           McNemar's Test P-Value : 0.00337
##
##           Sensitivity : 0.9918
##           Specificity : 0.8000
##           Pos Pred Value : 0.9773
##           Neg Pred Value : 0.9185
##           Prevalence : 0.8967
##           Detection Rate : 0.8893
##           Detection Prevalence : 0.9100
##           Balanced Accuracy : 0.8959
##
##           'Positive' Class : No
##
```

```

results.poly.both2 = prediction(pred.poly.both2, test_label)
roc.poly.both2 = performance(results.poly.both2, measure = "tpr", x.measure = "fpr")

model.poly.both3 = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + poly(CCAvg, 3) + Education + Securiti
es.Account + CD.Account + Online + CreditCard, family = "binomial", data = train)
pred.poly.both3 = predict(model.poly.both3, test, type = "response")
poly.both3.cutoff = 0.55
class.poly.both3 = as.factor(if_else(pred.poly.both3 < poly.both3.cutoff, "No", "Yes"))
prop.table(table(class.poly.both3))

```

```

## class.poly.both3
##           No           Yes
## 0.91066667 0.08933333

```

```

confusionMatrix(class.poly.both3, test$Personal.Loan)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1334  32
##           Yes  11 123
##
##           Accuracy : 0.9713
##           95% CI : (0.9616, 0.9792)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8354
##
##           Mcnemar's Test P-Value : 0.002289
##
##           Sensitivity : 0.9918
##           Specificity : 0.7935
##           Pos Pred Value : 0.9766
##           Neg Pred Value : 0.9179
##           Prevalence : 0.8967
##           Detection Rate : 0.8893
##           Detection Prevalence : 0.9107
##           Balanced Accuracy : 0.8927
##
##           'Positive' Class : No
##

```

```

results.poly.both3 = prediction(pred.poly.both3, test_label)
roc.poly.both3 = performance(results.poly.both3, measure = "tpr", x.measure = "fpr")

confusionMatrix(class.poly.income2, test$Personal.Loan)$overall[1]

```

```

## Accuracy
## 0.9706667

```

```

confusionMatrix(class.poly.income3, test$Personal.Loan)$overall[1]

```

```

## Accuracy
## 0.97

```

```

confusionMatrix(class.poly.CCAvg2, test$Personal.Loan)$overall[1]

```

```

## Accuracy
## 0.9673333

```

```

confusionMatrix(class.poly.CCAvg3, test$Personal.Loan)$overall[1]

```

```

## Accuracy
## 0.9666667

```

```

confusionMatrix(class.poly.both2, test$Personal.Loan)$overall[1]


```

```
## Accuracy
## 0.972
```

```
confusionMatrix(class.poly.both3, test$Personal.Loan)$overall[1]
```

```
## Accuracy
## 0.9713333
```

```
plot(roc.step, col = "red", xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.income2, col = "green", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.income3, col = "blue", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.CCAvg2, col = "blueviolet", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.CCAvg3, col = "aquamarine", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.both2, col = "chartreuse", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.both3, col = "coral", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
legend("bottomright", legend = c("Stepwise", "polyincome2", "polyincome3", "polyCCAVg2", "polyCCAVg3", "polyboth2",
", "polyboth3"),
      col = c("red", "green", "blue", "blueviolet", "aquamarine", "chartreuse", "coral"),
      lty=1, lwd=1)
```



Polynomial Income³ seems to be the best performing model according to the above plot.

Objective 2: Working Towards A Best Model

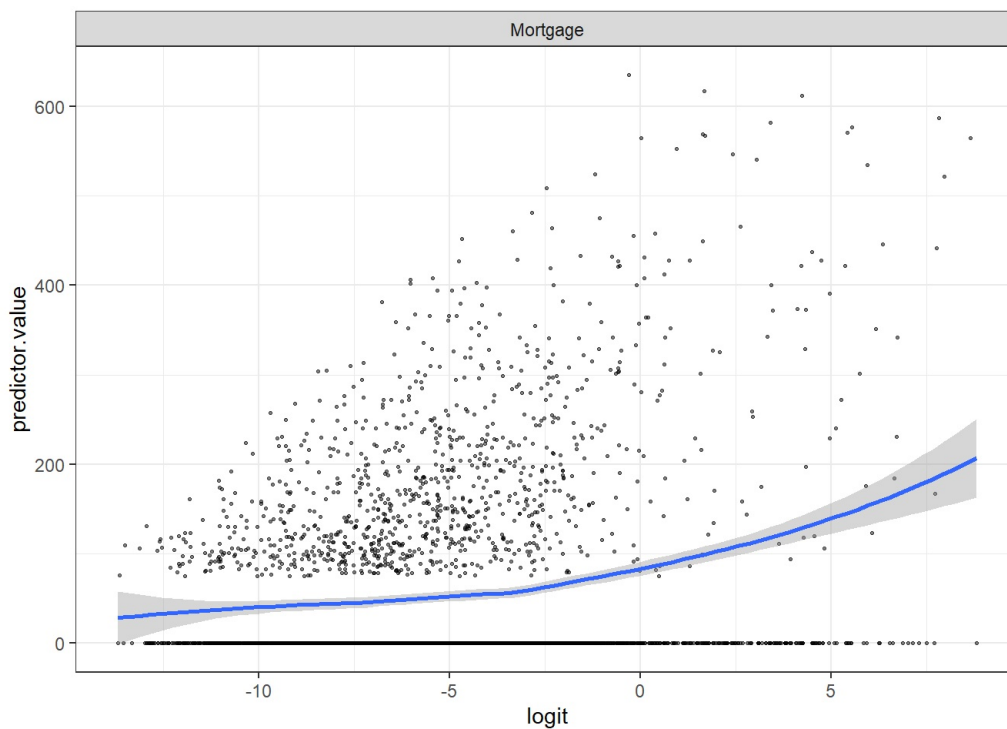
```
str(train)
```

```
## 'data.frame': 3500 obs. of 11 variables:
## $ Income : int 23 52 98 79 95 63 91 143 59 38 ...
## $ Family : Factor w/ 4 levels "1","2","3","4": 3 4 1 2 2 4 1 3 3 1 ...
## $ CCAvg : num 0.4 1.3 5.4 2.8 0 3.6 0.1 2.9 0.9 1.5 ...
## $ Education : Factor w/ 3 levels "1","2","3": 1 2 1 1 3 3 2 3 3 2 ...
## $ Mortgage : int 0 0 0 179 0 0 199 0 199 116 ...
## $ Personal.Loan : Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 2 1 1 ...
## $ Securities.Account: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 2 1 ...
## $ CD.Account : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ Online : Factor w/ 2 levels "0","1": 2 2 2 1 2 1 2 2 2 1 ...
## $ CreditCard : Factor w/ 2 levels "0","1": 1 2 1 1 1 1 2 1 1 2 ...
## $ Experience2 : num 0.538 0.613 0.04 0.594 0.636 ...
```



```
## Addressing Mortgage
facets = c("Mortgage")
ggplot(mydata[mydata$predictors %in% facets,], aes(logit, predictor.value))+
  geom_point(size = 0.5, alpha = 0.5) +
  geom_smooth(method = "loess") +
  theme_bw() +
  facet_wrap(~predictors, scales = "free_y")
```


```
## `geom_smooth()` using formula 'y ~ x'
```



```
## Training Stepwise again but with logged Mortgage
str(df$Mortgage)
```

```
## int [1:5000] 0 0 0 0 0 155 0 0 104 0 ...
```

```
final_df = mutate(df)
## Fixing Mortgage values
final_df$Mortgage[final_df$Mortgage == 0] = 1
final_df$LoggedMortgage = log(final_df$Mortgage)
plot(Personal.Loan ~ LoggedMortgage, data = final_df, col= c("pink","lightblue")) # Mortgage is now somewhat distributed
```



```
set.seed(123)
split = sample(nrow(final_df), nrow(final_df)*0.7)
final_train = final_df[split,]
final_test = final_df[-split,]
```

```
# prefinalmodel0: Stepwise with logged Mortgage.
prefinalmodel0 = glm(Personal.Loan ~ ., data = final_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefinalmodel0, direction = "both")
```

```
## Start: AIC=820.61
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + Experience2 +
##   LoggedMortgage
##
##           Df Deviance    AIC
## - Experience2      1   790.62  818.62
## - Mortgage         1   791.11  819.11
## - LoggedMortgage   1   792.42  820.42
## <none>              790.61  820.61
## - Securities.Account 1   795.74  823.74
## - CCAvg             1   795.92  823.92
## - CreditCard        1   803.52  831.52
## - Online            1   807.83  835.83
## - Family            3   864.66  888.66
## - CD.Account        1   878.76  906.76
## - Education         2  1082.26 1108.26
## - Income            1  1333.24 1361.24
##
## Step: AIC=818.62
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## - Mortgage         1   791.11  817.11
## - LoggedMortgage   1   792.43  818.43
## <none>              790.62  818.62
## + Experience2      1   790.61  820.61
## - Securities.Account 1   795.74  821.74
## - CCAvg            1   796.01  822.01
## - CreditCard       1   803.52  829.52
## - Online           1   807.84  833.84
## - Family           3   864.66  886.66
## - CD.Account       1   878.95  904.95
## - Education        2  1082.32 1106.32
## - Income           1  1333.25 1359.25
##
## Step: AIC=817.11
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
##   CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## <none>              791.11  817.11
## - LoggedMortgage   1   794.07  818.07
## + Mortgage         1   790.62  818.62
## + Experience2      1   791.11  819.11
## - Securities.Account 1   796.05  820.05
## - CCAvg            1   796.93  820.93
## - CreditCard       1   803.83  827.83
## - Online           1   808.18  832.18
## - Family           3   864.83  884.83
## - CD.Account       1   879.01  903.01
## - Education        2  1082.53 1104.53
## - Income           1  1381.04 1405.04
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
##   family = "binomial", data = final_train)
##
## Coefficients:
##      (Intercept)              Income              Family2
##      -12.64983              0.06365              -0.24090
##      Family3              Family4              CCAvg
##      1.91321              1.38231              0.13262
##      Education2              Education3  Securities.Account1
##      4.03980              4.20467              -0.74697
##      CD.Account1              Online1              CreditCard1
##      3.52663              -0.82814              -0.88255
##      LoggedMortgage
##      0.06474
##
## Degrees of Freedom: 3499 Total (i.e. Null); 3487 Residual
## Null Deviance: 2164
## Residual Deviance: 791.1    AIC: 817.1
```

```
finalmodel0 = glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
  Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
  family = "binomial", data = final_train)
summary(finalmodel0)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
##   family = "binomial", data = final_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8579  -0.1805  -0.0646  -0.0192   4.1186
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -12.649830   0.669463 -18.895 < 2e-16 ***
## Income         0.063648   0.003719  17.116 < 2e-16 ***
## Family2       -0.240898   0.282962  -0.851 0.394579
## Family3        1.913214   0.295328   6.478 9.28e-11 ***
## Family4        1.382308   0.290321   4.761 1.92e-06 ***
## CCAvg          0.132625   0.055363   2.396 0.016595 *
## Education2     4.039801   0.335595  12.038 < 2e-16 ***
## Education3     4.204674   0.335101  12.547 < 2e-16 ***
## Securities.Account1 -0.746968  0.350753  -2.130 0.033204 *
## CD.Account1     3.526629   0.403314   8.744 < 2e-16 ***
## Online1        -0.828144   0.202636  -4.087 4.37e-05 ***
## CreditCard1    -0.882548   0.259549  -3.400 0.000673 ***
## LoggedMortgage  0.064738   0.037410   1.730 0.083544 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2163.69  on 3499  degrees of freedom
## Residual deviance:  791.11  on 3487  degrees of freedom
## AIC: 817.11
##
## Number of Fisher Scoring iterations: 8
```

```
AIC(finalmodel0)      # AIC = 817.1093
```

```
## [1] 817.1093
```

```
BIC(finalmodel0)      # BIC = 897.1968
```

```
## [1] 897.196
```

```
## Testing finalmodel0
pred.final0 = predict(finalmodel0, final_test, type = "response")
final0.cutoff = 0.4
class.final0 = as.factor(if_else(pred.final0 < final0.cutoff, "No", "Yes"))
prop.table(table(class.final0))
```

```
## class.final0
##           No           Yes
## 0.91266667 0.08733333
```

```
## Confusion Matrix
confusionMatrix(class.final0, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1333  36
##           Yes  12 119
##
##           Accuracy : 0.968
##           95% CI : (0.9578, 0.9763)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8146
##
## Mcnemar's Test P-Value : 0.0009009
##
##           Sensitivity : 0.9911
##           Specificity : 0.7677
##           Pos Pred Value : 0.9737
##           Neg Pred Value : 0.9084
##           Prevalence : 0.8967
##           Detection Rate : 0.8887
##           Detection Prevalence : 0.9127
##           Balanced Accuracy : 0.8794
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.4
# Accuracy = 0.968
# Sensitivity = 0.76774
# Specificity = 0.99108
```

```
# prefinalmodel0a: Stepwise with logged Mortgage and 3rd power Income.
prefinalmodel0a = glm(Personal.Loan ~ ., data = final_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefinalmodel0a, direction = "both")
```

```
## Start: AIC=820.61
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + Experience2 +
##   LoggedMortgage
##
##           Df Deviance    AIC
## - Experience2      1   790.62  818.62
## - Mortgage         1   791.11  819.11
## - LoggedMortgage    1   792.42  820.42
## <none>              790.61  820.61
## - Securities.Account 1   795.74  823.74
## - CCAvg             1   795.92  823.92
## - CreditCard        1   803.52  831.52
## - Online            1   807.83  835.83
## - Family            3   864.66  888.66
## - CD.Account        1   878.76  906.76
## - Education         2  1082.26 1108.26
## - Income            1  1333.24 1361.24
##
## Step: AIC=818.62
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## - Mortgage         1   791.11  817.11
## - LoggedMortgage    1   792.43  818.43
## <none>              790.62  818.62
## + Experience2      1   790.61  820.61
## - Securities.Account 1   795.74  821.74
## - CCAvg            1   796.01  822.01
## - CreditCard       1   803.52  829.52
## - Online           1   807.84  833.84
## - Family           3   864.66  886.66
## - CD.Account       1   878.95  904.95
## - Education        2  1082.32 1106.32
## - Income           1  1333.25 1359.25
##
## Step: AIC=817.11
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
##   CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## <none>              791.11  817.11
## - LoggedMortgage    1   794.07  818.07
## + Mortgage         1   790.62  818.62
## + Experience2      1   791.11  819.11
## - Securities.Account 1   796.05  820.05
## - CCAvg            1   796.93  820.93
## - CreditCard       1   803.83  827.83
## - Online           1   808.18  832.18
## - Family           3   864.83  884.83
## - CD.Account       1   879.01  903.01
## - Education        2  1082.53 1104.53
## - Income           1  1381.04 1405.04
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
##   family = "binomial", data = final_train)
##
## Coefficients:
##      (Intercept)          Income          Family2
##      -12.64983         0.06365        -0.24090
##      Family3          Family4          CCAvg
##      1.91321         1.38231         0.13262
##      Education2      Education3  Securities.Account1
##      4.03980         4.20467        -0.74697
##      CD.Account1      Online1      CreditCard1
##      3.52663        -0.82814        -0.88255
##      LoggedMortgage
##      0.06474
##
## Degrees of Freedom: 3499 Total (i.e. Null); 3487 Residual
## Null Deviance:      2164
## Residual Deviance: 791.1    AIC: 817.1
```

```
finalmodel0a = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + CCAvg + Education +
  Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
  family = "binomial", data = final_train)
summary(finalmodel0a)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ poly(Income, 3) + Family + CCAvg +
##   Education + Securities.Account + CD.Account + Online + CreditCard +
##   LoggedMortgage, family = "binomial", data = final_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5046  -0.1067  -0.0085  -0.0004   4.7966
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -11.70723     1.13599  -10.306 < 2e-16 ***
## poly(Income, 3)1    357.28275    57.75651   6.186 6.17e-10 ***
## poly(Income, 3)2   -128.05811    41.07901  -3.117 0.001825 **
## poly(Income, 3)3    -4.65606    18.48375  -0.252 0.801118
## Family2          -0.17600     0.28292  -0.622 0.533884
## Family3           2.47418     0.33878   7.303 2.81e-13 ***
## Family4           1.59917     0.32235   4.961 7.01e-07 ***
## CCAvg             0.16062     0.05588   2.874 0.004048 **
## Education2        3.88617     0.32583  11.927 < 2e-16 ***
## Education3        4.02449     0.32635  12.332 < 2e-16 ***
## Securities.Account1 -0.85671     0.40702  -2.105 0.035304 *
## CD.Account1       3.75879     0.47962   7.837 4.61e-15 ***
## Online1          -0.86468     0.22927  -3.771 0.000162 ***
## CreditCard1      -0.96375     0.29073  -3.315 0.000917 ***
## LoggedMortgage     0.09663     0.04152   2.327 0.019958 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2163.7  on 3499  degrees of freedom
## Residual deviance:  625.5  on 3485  degrees of freedom
## AIC: 655.5
##
## Number of Fisher Scoring iterations: 11
```

```
AIC(finalmodel0a)      # AIC = 655.5022
```

```
## [1] 655.5022
```

```
BIC(finalmodel0a)      # BIC = 747.9099
```

```
## [1] 747.9099
```

```
## Testing finalmodela
pred.final0a = predict(finalmodel0a, final_test, type = "response")
final0a.cutoff = 0.33
class.final0a = as.factor(if_else(pred.final0a < final0a.cutoff, "No", "Yes"))
prop.table(table(class.final0a))
```

```
## class.final0a
##   No  Yes
## 0.89 0.11
```

```
## Confusion Matrix
confusionMatrix(class.final0a, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1319  16
##           Yes  26 139
##
##           Accuracy : 0.972
##           95% CI : (0.9623, 0.9797)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.8531
##
## Mcnemar's Test P-Value : 0.1649
##
##           Sensitivity : 0.9807
##           Specificity : 0.8968
##           Pos Pred Value : 0.9880
##           Neg Pred Value : 0.8424
##           Prevalence : 0.8967
##           Detection Rate : 0.8793
##           Detection Prevalence : 0.8900
##           Balanced Accuracy : 0.9387
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.33
# Accuracy = 0.972
# Sensitivity = 0.89677
# Specificity = 0.98067
```

```
# prefinalmodell1: Stepwise with logged Mortgage and below observations removed.
modell_train = final_train[-c(350, 976, 1070, 1127, 2159),]
prefinalmodell1 = glm(Personal.Loan ~ ., data = modell_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefinalmodell1, direction = "both")
```



```
## Start: AIC=817.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + Experience2 +
##   LoggedMortgage
##
##           Df Deviance    AIC
## - Experience2      1   787.42  815.42
## - Mortgage         1   787.94  815.94
## - LoggedMortgage    1   789.22  817.22
## <none>              787.42  817.42
## - Securities.Account 1   792.70  820.70
## - CCAvg             1   793.53  821.53
## - CreditCard        1   800.65  828.65
## - Online            1   805.10  833.10
## - Family            3   861.44  885.44
## - CD.Account        1   875.43  903.43
## - Education         2  1077.58 1103.58
## - Income            1  1331.23 1359.23
##
## Step: AIC=815.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## - Mortgage         1   787.94  813.94
## - LoggedMortgage    1   789.23  815.23
## <none>              787.42  815.42
## + Experience2      1   787.42  817.42
## - Securities.Account 1   792.70  818.70
## - CCAvg             1   793.60  819.60
## - CreditCard        1   800.66  826.66
## - Online            1   805.10  831.10
## - Family            3   861.44  883.44
## - CD.Account        1   875.67  901.67
## - Education         2  1077.63 1101.63
## - Income            1  1331.24 1357.24
##
## Step: AIC=813.94
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
##   CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## <none>              787.94  813.94
## - LoggedMortgage    1   790.73  814.73
## + Mortgage         1   787.42  815.42
## + Experience2      1   787.94  815.94
## - Securities.Account 1   793.03  817.03
## - CCAvg             1   794.58  818.58
## - CreditCard        1   800.99  824.99
## - Online            1   805.46  829.46
## - Family            3   861.63  881.63
## - CD.Account        1   875.74  899.74
## - Education         2  1077.85 1099.85
## - Income            1  1378.66 1402.66
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
##   family = "binomial", data = modell_train)
##
## Coefficients:
##      (Intercept)              Income              Family2
##      -12.63965              0.06374             -0.28917
##      Family3              Family4              CCAvg
##      1.88835              1.35862              0.14255
##      Education2          Education3  Securities.Account1
##      4.03491              4.19975             -0.75893
##      CD.Account1          Online1              CreditCard1
##      3.53185             -0.84114             -0.89448
##      LoggedMortgage
##      0.06285
##
## Degrees of Freedom: 3494 Total (i.e. Null); 3482 Residual
## Null Deviance: 2163
## Residual Deviance: 787.9    AIC: 813.9
```

```
finalmodell = glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +  
  Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,  
  family = "binomial", data = modell_train)  
summary(finalmodell)
```

```
##  
## Call:  
## glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +  
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,  
##   family = "binomial", data = modell_train)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max  
## -2.8681  -0.1811  -0.0640  -0.0188   4.1208  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)   -12.639653   0.669128 -18.890 < 2e-16 ***  
## Income         0.063743   0.003725  17.114 < 2e-16 ***  
## Family2       -0.289167   0.284194  -1.017 0.308917  
## Family3        1.888351   0.295229   6.396 1.59e-10 ***  
## Family4        1.358621   0.290154   4.682 2.84e-06 ***  
## CCAvg          0.142553   0.055765   2.556 0.010578 *  
## Education2     4.034915   0.335867  12.013 < 2e-16 ***  
## Education3     4.199750   0.335388  12.522 < 2e-16 ***  
## Securities.Account1 -0.758926  0.351432  -2.160 0.030810 *  
## CD.Account1    3.531847   0.404355   8.735 < 2e-16 ***  
## Online1       -0.841139   0.203271  -4.138 3.50e-05 ***  
## CreditCard1   -0.894477   0.259920  -3.441 0.000579 ***  
## LoggedMortgage  0.062850   0.037466   1.678 0.093442 .  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##    Null deviance: 2162.72  on 3494  degrees of freedom  
## Residual deviance: 787.94  on 3482  degrees of freedom  
## AIC: 813.94  
##  
## Number of Fisher Scoring iterations: 8
```

```
AIC(finalmodell)      # AIC = 813.9434
```

```
## [1] 813.9434
```

```
BIC(finalmodell)      # BIC = 894.0116
```

```
## [1] 894.0116
```

```
## Testing finalmodell  
pred.final1 = predict(finalmodell, final_test, type = "response")  
final1.cutoff = 0.4  
class.final1 = as.factor(if_else(pred.final1 < final1.cutoff, "No", "Yes"))  
prop.table(table(class.final1))
```

```
## class.final1  
##           No           Yes  
## 0.91066667 0.08933333
```

```
## Confusion Matrix  
confusionMatrix(class.final1, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1331  35
##           Yes  14 120
##
##           Accuracy : 0.9673
##           95% CI : (0.957, 0.9757)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8125
##
## Mcnemar's Test P-Value : 0.004275
##
##           Sensitivity : 0.9896
##           Specificity : 0.7742
##           Pos Pred Value : 0.9744
##           Neg Pred Value : 0.8955
##           Prevalence : 0.8967
##           Detection Rate : 0.8873
##           Detection Prevalence : 0.9107
##           Balanced Accuracy : 0.8819
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.4
# Accuracy = 0.9673
# Sensitivity = 0.77419
# Specificity = 0.98959
```

```
# prefinalmodella: Stepwise with logged Mortgage, below observations removed, and 3rd power Income.
modell_train = final_train[-c(350, 976, 1070, 1127, 2159),]
prefinalmodella = glm(Personal.Loan ~ ., data = modell_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefinalmodella, direction = "both")
```

```
## Start: AIC=817.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
## Securities.Account + CD.Account + Online + CreditCard + Experience2 +
## LoggedMortgage
##
## Df Deviance AIC
## - Experience2 1 787.42 815.42
## - Mortgage 1 787.94 815.94
## - LoggedMortgage 1 789.22 817.22
## <none> 787.42 817.42
## - Securities.Account 1 792.70 820.70
## - CCAvg 1 793.53 821.53
## - CreditCard 1 800.65 828.65
## - Online 1 805.10 833.10
## - Family 3 861.44 885.44
## - CD.Account 1 875.43 903.43
## - Education 2 1077.58 1103.58
## - Income 1 1331.23 1359.23
##
## Step: AIC=815.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
## Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
## Df Deviance AIC
## - Mortgage 1 787.94 813.94
## - LoggedMortgage 1 789.23 815.23
## <none> 787.42 815.42
## + Experience2 1 787.42 817.42
## - Securities.Account 1 792.70 818.70
## - CCAvg 1 793.60 819.60
## - CreditCard 1 800.66 826.66
## - Online 1 805.10 831.10
## - Family 3 861.44 883.44
## - CD.Account 1 875.67 901.67
## - Education 2 1077.63 1101.63
## - Income 1 1331.24 1357.24
##
## Step: AIC=813.94
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
## CD.Account + Online + CreditCard + LoggedMortgage
##
## Df Deviance AIC
## <none> 787.94 813.94
## - LoggedMortgage 1 790.73 814.73
## + Mortgage 1 787.42 815.42
## + Experience2 1 787.94 815.94
## - Securities.Account 1 793.03 817.03
## - CCAvg 1 794.58 818.58
## - CreditCard 1 800.99 824.99
## - Online 1 805.46 829.46
## - Family 3 861.63 881.63
## - CD.Account 1 875.74 899.74
## - Education 2 1077.85 1099.85
## - Income 1 1378.66 1402.66
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
## Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
## family = "binomial", data = modell_train)
##
## Coefficients:
## (Intercept) Income Family2
## -12.63965 0.06374 -0.28917
## Family3 Family4 CCAvg
## 1.88835 1.35862 0.14255
## Education2 Education3 Securities.Account1
## 4.03491 4.19975 -0.75893
## CD.Account1 Online1 CreditCard1
## 3.53185 -0.84114 -0.89448
## LoggedMortgage
## 0.06285
##
## Degrees of Freedom: 3494 Total (i.e. Null); 3482 Residual
## Null Deviance: 2163
## Residual Deviance: 787.9 AIC: 813.9
```

```
finalmodella = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + CCAvg + Education +  
  Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,  
  family = "binomial", data = modell_train)  
summary(finalmodella)
```

```
##  
## Call:  
## glm(formula = Personal.Loan ~ poly(Income, 3) + Family + CCAvg +  
##   Education + Securities.Account + CD.Account + Online + CreditCard +  
##   LoggedMortgage, family = "binomial", data = modell_train)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max  
## -2.4994  -0.1072  -0.0085  -0.0004   4.7967  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)    -11.71235     1.13606  -10.310 < 2e-16 ***  
## poly(Income, 3)1    357.74419    57.71361   6.199 5.70e-10 ***  
## poly(Income, 3)2   -128.16066    40.98097  -3.127 0.001764 **  
## poly(Income, 3)3    -4.09516    18.41518  -0.222 0.824018  
## Family2          -0.18928     0.28343  -0.668 0.504250  
## Family3           2.46348     0.33886   7.270 3.60e-13 ***  
## Family4           1.59104     0.32230   4.936 7.96e-07 ***  
## CCAvg             0.16322     0.05603   2.913 0.003578 **  
## Education2         3.88120     0.32577  11.914 < 2e-16 ***  
## Education3         4.02008     0.32633  12.319 < 2e-16 ***  
## Securities.Account1 -0.85804     0.40683  -2.109 0.034937 *  
## CD.Account1        3.75405     0.47939   7.831 4.84e-15 ***  
## Online1           -0.86466     0.22930  -3.771 0.000163 ***  
## CreditCard1        -0.96647     0.29056  -3.326 0.000880 ***  
## LoggedMortgage      0.09568     0.04151   2.305 0.021173 *  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 2162.7  on 3494  degrees of freedom  
## Residual deviance:  624.9  on 3480  degrees of freedom  
## AIC: 654.9  
##  
## Number of Fisher Scoring iterations: 11
```

```
AIC(finalmodella)      # AIC = 654.8965
```

```
## [1] 654.8965
```

```
BIC(finalmodella)      # BIC = 747.2828
```

```
## [1] 747.2828
```

```
## Testing finalmodella  
pred.finalla = predict(finalmodella, final_test, type = "response")  
finalla.cutoff = 0.33  
class.finalla = as.factor(if_else(pred.finalla < finalla.cutoff, "No", "Yes"))  
prop.table(table(class.finalla))
```

```
## class.finalla  
##    No  Yes  
## 0.89 0.11
```

```
## Confusion Matrix  
confusionMatrix(class.finalla, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 1319  16
##           Yes  26 139
##
##           Accuracy : 0.972
##           95% CI : (0.9623, 0.9797)
##           No Information Rate : 0.8967
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.8531
##
## Mcnemar's Test P-Value : 0.1649
##
##           Sensitivity : 0.9807
##           Specificity : 0.8968
##           Pos Pred Value : 0.9880
##           Neg Pred Value : 0.8424
##           Prevalence : 0.8967
##           Detection Rate : 0.8793
##           Detection Prevalence : 0.8900
##           Balanced Accuracy : 0.9387
##
##           'Positive' Class : No
##
```

```
# Threshold = 0.33
# Accuracy = 0.972
# Sensitivity = 0.89677
# Specificity = 0.98067
```

Objective 2: Inspecting Interactions

```
# This is using the original dataset
```

```
library(sjPlot) #For effect plotting
```

```
## Warning: package 'sjPlot' was built under R version 4.1.3
```

```
## Learn more about sjPlot with 'browseVignettes("sjPlot")'.
```

```
library(sjmisc) #For effect plotting
```

```
## Warning: package 'sjmisc' was built under R version 4.1.3
```

```
##
## Attaching package: 'sjmisc'
```

```
## The following objects are masked from 'package:jtools':
##
## %nin%, center
```

```
## The following object is masked from 'package:purrr':
##
## is_empty
```

```
## The following object is masked from 'package:tidyr':
##
## replace_na
```

```
## The following object is masked from 'package:tibble':
##
## add_case
```

```
getwd()
```

```
## [1] "C:/Users/dnguy/Desktop/2 Applied Stats/Project 2/Statistics2-project-2"
```

```
Personall=read.csv("Bank_Personal_Loan_Modelling.csv")

# Age is omitted for plotting since it's replaced by Experience2.

#Education,Family,CCAvg,Online,CreditCard,Securities.Account
a=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=Education))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Education)

b0=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=Family))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Family)

c=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=CCAvg))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CCAvg)

d=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=Online))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Online)

e=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=CreditCard))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CreditCard)

f=ggplot(Personall,aes(x=Income,y=Personal.Loan,colour=Securities.Account))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Securities.Account)

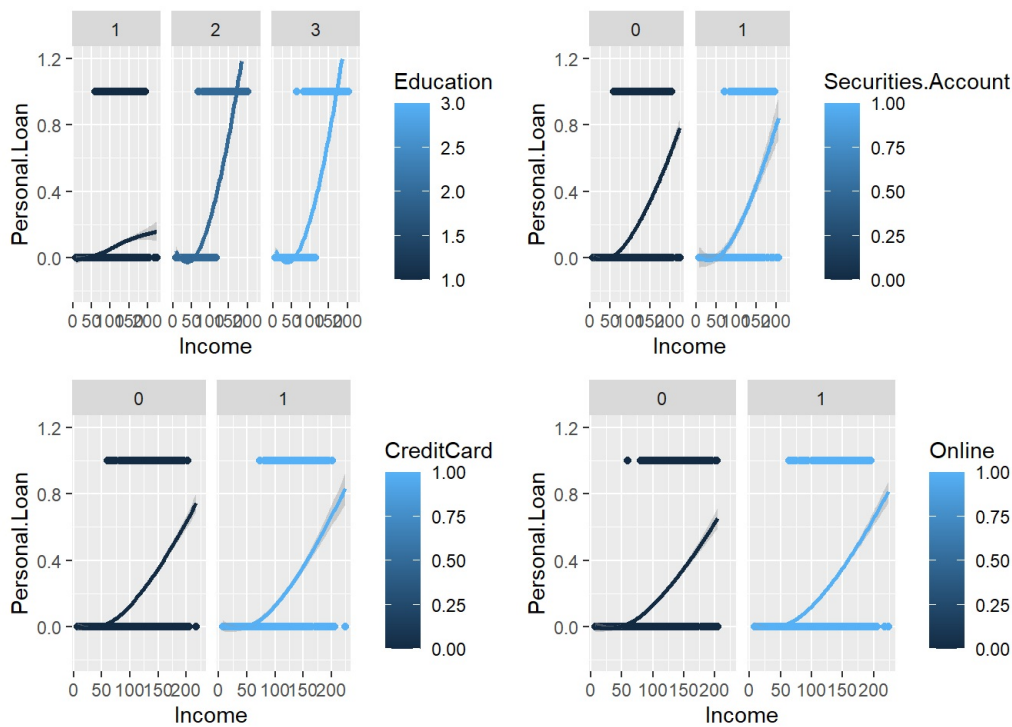
grid.arrange(a,f,e,d, ncol=2)
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
## Warning: Removed 14 rows containing missing values (geom_smooth).
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```



```
#Education,Family,CCAvg,Online,CreditCard, Securities.Account
a=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=Education))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Education)

b1=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=Family))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Family)

c=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=CCAvg))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CCAvg)

d=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=Online))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Online)

e=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=CreditCard))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CreditCard)

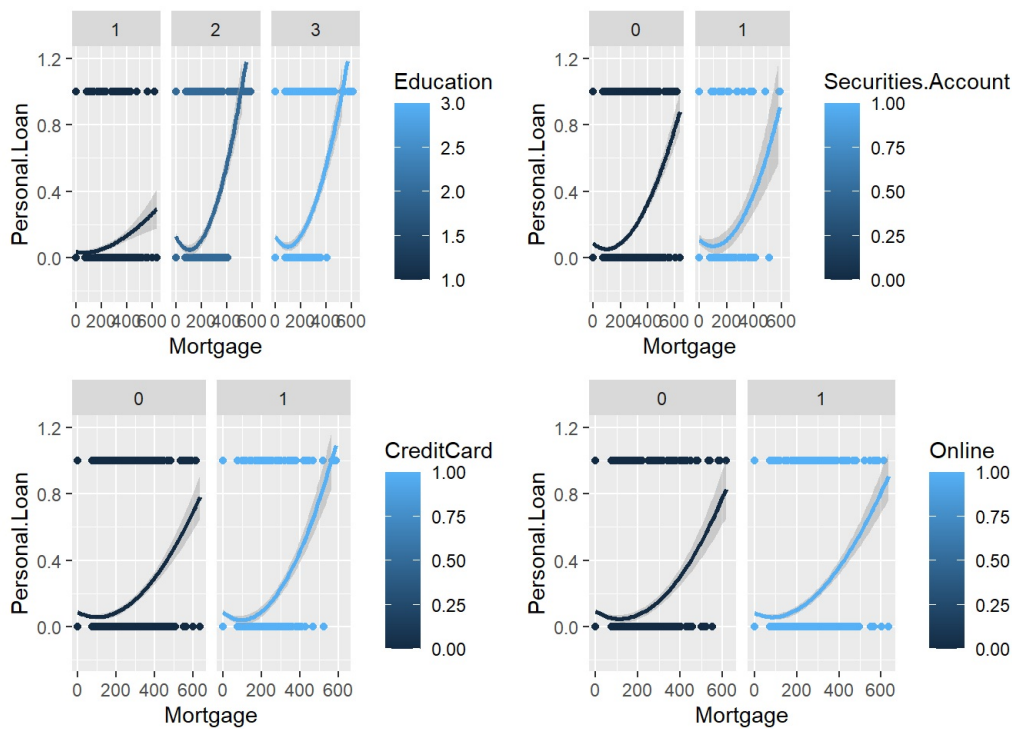
f=ggplot(PersonalL,aes(x=Mortgage,y=Personal.Loan,colour=Securities.Account))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Securities.Account)

grid.arrange(a,f,e,d, ncol=2)
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
## Warning: Removed 11 rows containing missing values (geom_smooth).
```

```
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```

```
#Education,Family,CCAvg,Online,CreditCard, Securities.Account
a=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=Education))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Education)

b2=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=Family))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Family)

c=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=CCAvg))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CCAvg)

d=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=Online))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Online)

e=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=CreditCard))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~CreditCard)

f=ggplot(PersonalL,aes(x=CCAvg,y=Personal.Loan,colour=Securities.Account))+geom_point()+
  geom_smooth(method="loess",size=1,span=1.5)+
  ylim(-.2,1.2)+
  facet_wrap(~Securities.Account)

# prefinalmodel2: Stepwise with logged Mortgage, below observations removed, 3rd power Income, and Income*Family.
modell_train = final_train[-c(350, 976, 1070, 1127, 2159),]
prefinalmodel2 = glm(Personal.Loan ~ ., data = modell_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefinalmodel2, direction = "both")
```

```
## Start: AIC=817.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
## Securities.Account + CD.Account + Online + CreditCard + Experience2 +
## LoggedMortgage
##
## Df Deviance AIC
## - Experience2 1 787.42 815.42
## - Mortgage 1 787.94 815.94
## - LoggedMortgage 1 789.22 817.22
## <none> 787.42 817.42
## - Securities.Account 1 792.70 820.70
## - CCAvg 1 793.53 821.53
## - CreditCard 1 800.65 828.65
## - Online 1 805.10 833.10
## - Family 3 861.44 885.44
## - CD.Account 1 875.43 903.43
## - Education 2 1077.58 1103.58
## - Income 1 1331.23 1359.23
##
## Step: AIC=815.42
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
## Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
## Df Deviance AIC
## - Mortgage 1 787.94 813.94
## - LoggedMortgage 1 789.23 815.23
## <none> 787.42 815.42
## + Experience2 1 787.42 817.42
## - Securities.Account 1 792.70 818.70
## - CCAvg 1 793.60 819.60
## - CreditCard 1 800.66 826.66
## - Online 1 805.10 831.10
## - Family 3 861.44 883.44
## - CD.Account 1 875.67 901.67
## - Education 2 1077.63 1101.63
## - Income 1 1331.24 1357.24
##
## Step: AIC=813.94
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
## CD.Account + Online + CreditCard + LoggedMortgage
##
## Df Deviance AIC
## <none> 787.94 813.94
## - LoggedMortgage 1 790.73 814.73
## + Mortgage 1 787.42 815.42
## + Experience2 1 787.94 815.94
## - Securities.Account 1 793.03 817.03
## - CCAvg 1 794.58 818.58
## - CreditCard 1 800.99 824.99
## - Online 1 805.46 829.46
## - Family 3 861.63 881.63
## - CD.Account 1 875.74 899.74
## - Education 2 1077.85 1099.85
## - Income 1 1378.66 1402.66
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
## Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
## family = "binomial", data = modell_train)
##
## Coefficients:
## (Intercept) Income Family2
## -12.63965 0.06374 -0.28917
## Family3 Family4 CCAvg
## 1.88835 1.35862 0.14255
## Education2 Education3 Securities.Account1
## 4.03491 4.19975 -0.75893
## CD.Account1 Online1 CreditCard1
## 3.53185 -0.84114 -0.89448
## LoggedMortgage
## 0.06285
##
## Degrees of Freedom: 3494 Total (i.e. Null); 3482 Residual
## Null Deviance: 2163
## Residual Deviance: 787.9 AIC: 813.9
```

```
finalmodel2 = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + Income*Family + CCAvg + Education + Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
  family = "binomial", data = modell_train)
#finalmodel2 = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + Income*Family + CCAvg + CCAvg*Family + Education + CCAvg*Education + Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage + LoggedMortgage*Education + LoggedMortgage*Family, family = "binomial", data = modell_train)
summary(finalmodel2)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ poly(Income, 3) + Family + Income *
##   Family + CCAvg + Education + Securities.Account + CD.Account +
##   Online + CreditCard + LoggedMortgage, family = "binomial",
##   data = modell_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4936  -0.1100  -0.0175  -0.0015   4.2247
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -8.684063    0.885151  -9.811 < 2e-16 ***
## poly(Income, 3)1  193.117742   44.684965   4.322 1.55e-05 ***
## poly(Income, 3)2  -53.501216   32.211896  -1.661 0.096731 .
## poly(Income, 3)3  -19.228152   16.524968  -1.164 0.244594
## Family2          0.104771    1.053352   0.099 0.920770
## Family3         -0.076632    2.321404  -0.329 0.000503 ***
## Family4        -10.301547    2.401639  -4.289 1.79e-05 ***
## Income              NA          NA      NA      NA
## CCAvg              0.232488    0.064788   3.588 0.000333 ***
## Education2         4.293987    0.359126  11.957 < 2e-16 ***
## Education3         4.271829    0.352929  12.104 < 2e-16 ***
## Securities.Account1 -0.824892    0.428669  -1.924 0.054316 .
## CD.Account1        3.571350    0.507698   7.034 2.00e-12 ***
## Online1           -0.863168    0.251394  -3.434 0.000596 ***
## CreditCard1       -1.027839    0.311468  -3.300 0.000967 ***
## LoggedMortgage     0.050310    0.045351   1.109 0.267277
## Family2:Income     -0.001187    0.007887  -0.151 0.880340
## Family3:Income     0.092702    0.021798   4.253 2.11e-05 ***
## Family4:Income     0.106746    0.021944   4.864 1.15e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 2162.72  on 3494  degrees of freedom
## Residual deviance: 548.51  on 3477  degrees of freedom
## AIC: 584.51
##
## Number of Fisher Scoring iterations: 11
```

```
AIC(finalmodel2)      # AIC = 549.3415
```

```
## [1] 584.5078
```

```
BIC(finalmodel2)      # BIC = 678.6824
```

```
## [1] 695.3714
```

```
## Testing finalmodel2
pred.final2 = predict(finalmodel2, final_test, type = "response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
final2.cutoff = 0.34
class.final2 = as.factor(if_else(pred.final2 < final2.cutoff, "No", "Yes"))
prop.table(table(class.final2))
```

```
## class.final2
##      No      Yes
## 0.8906667 0.1093333
```

```
## Confusion Matrix
confusionMatrix(class.final2, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  No  Yes
##      No 1321  15
##      Yes  24 140
##
##      Accuracy : 0.974
##      95% CI : (0.9646, 0.9814)
##      No Information Rate : 0.8967
##      P-Value [Acc > NIR] : <2e-16
##
##      Kappa : 0.8632
##
##      Mcnemar's Test P-Value : 0.2002
##
##      Sensitivity : 0.9822
##      Specificity : 0.9032
##      Pos Pred Value : 0.9888
##      Neg Pred Value : 0.8537
##      Prevalence : 0.8967
##      Detection Rate : 0.8807
##      Detection Prevalence : 0.8907
##      Balanced Accuracy : 0.9427
##
##      'Positive' Class : No
##
```

```
# Threshold = 0.34
# Accuracy = 0.974
# Sensitivity = 0.90323
# Specificity = 0.98216
```

```
# prefina1model2a: Stepwise with logged Mortgage, 3rd power Income, and Income*Family.
prefina1model2a = glm(Personal.Loan ~ ., data = final_train, family = "binomial")
## Feature selection - stepwise
stepAIC(prefina1model2a, direction = "both")
```

```
## Start: AIC=820.61
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + Experience2 +
##   LoggedMortgage
##
##           Df Deviance    AIC
## - Experience2      1   790.62  818.62
## - Mortgage         1   791.11  819.11
## - LoggedMortgage   1   792.42  820.42
## <none>              790.61  820.61
## - Securities.Account 1   795.74  823.74
## - CCAvg             1   795.92  823.92
## - CreditCard        1   803.52  831.52
## - Online            1   807.83  835.83
## - Family            3   864.66  888.66
## - CD.Account        1   878.76  906.76
## - Education         2  1082.26 1108.26
## - Income            1  1333.24 1361.24
##
## Step: AIC=818.62
## Personal.Loan ~ Income + Family + CCAvg + Education + Mortgage +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## - Mortgage         1   791.11  817.11
## - LoggedMortgage   1   792.43  818.43
## <none>              790.62  818.62
## + Experience2      1   790.61  820.61
## - Securities.Account 1   795.74  821.74
## - CCAvg             1   796.01  822.01
## - CreditCard        1   803.52  829.52
## - Online            1   807.84  833.84
## - Family            3   864.66  886.66
## - CD.Account        1   878.95  904.95
## - Education         2  1082.32 1106.32
## - Income            1  1333.25 1359.25
##
## Step: AIC=817.11
## Personal.Loan ~ Income + Family + CCAvg + Education + Securities.Account +
##   CD.Account + Online + CreditCard + LoggedMortgage
##
##           Df Deviance    AIC
## <none>              791.11  817.11
## - LoggedMortgage   1   794.07  818.07
## + Mortgage         1   790.62  818.62
## + Experience2      1   791.11  819.11
## - Securities.Account 1   796.05  820.05
## - CCAvg             1   796.93  820.93
## - CreditCard        1   803.83  827.83
## - Online            1   808.18  832.18
## - Family            3   864.83  884.83
## - CD.Account        1   879.01  903.01
## - Education         2  1082.53 1104.53
## - Income            1  1381.04 1405.04
```

```
##
## Call: glm(formula = Personal.Loan ~ Income + Family + CCAvg + Education +
##   Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
##   family = "binomial", data = final_train)
##
## Coefficients:
##      (Intercept)              Income              Family2
##      -12.64983              0.06365              -0.24090
##      Family3              Family4              CCAvg
##      1.91321              1.38231              0.13262
##      Education2              Education3  Securities.Account1
##      4.03980              4.20467              -0.74697
##      CD.Account1              Online1              CreditCard1
##      3.52663              -0.82814              -0.88255
##      LoggedMortgage
##      0.06474
##
## Degrees of Freedom: 3499 Total (i.e. Null); 3487 Residual
## Null Deviance: 2164
## Residual Deviance: 791.1    AIC: 817.1
```

```
finalmodel2a = glm(formula = Personal.Loan ~ poly(Income, 3) + Family + Income*Family + CCAvg + Education + Securities.Account + CD.Account + Online + CreditCard + LoggedMortgage,
  family = "binomial", data = modell_train)
```

```
summary(finalmodel2a)
```

```
##
## Call:
## glm(formula = Personal.Loan ~ poly(Income, 3) + Family + Income *
##     Family + CCAvg + Education + Securities.Account + CD.Account +
##     Online + CreditCard + LoggedMortgage, family = "binomial",
##     data = modell_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4936  -0.1100  -0.0175  -0.0015   4.2247
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -8.684063    0.885151  -9.811 < 2e-16 ***
## poly(Income, 3)1  193.117742   44.684965   4.322 1.55e-05 ***
## poly(Income, 3)2  -53.501216   32.211896  -1.661 0.096731 .
## poly(Income, 3)3  -19.228152   16.524968  -1.164 0.244594
## Family2          0.104771    1.053352   0.099 0.920770
## Family3         -8.076632    2.321404  -3.479 0.000503 ***
## Family4        -10.301547    2.401639  -4.289 1.79e-05 ***
## Income              NA              NA      NA      NA
## CCAvg             0.232488    0.064788   3.588 0.000333 ***
## Education2        4.293987    0.359126  11.957 < 2e-16 ***
## Education3        4.271829    0.352929  12.104 < 2e-16 ***
## Securities.Account1 -0.824892    0.428669  -1.924 0.054316 .
## CD.Account1       3.571350    0.507698   7.034 2.00e-12 ***
## Online1          -0.863168    0.251394  -3.434 0.000596 ***
## CreditCard1      -1.027839    0.311468  -3.300 0.000967 ***
## LoggedMortgage     0.050310    0.045351   1.109 0.267277
## Family2:Income    -0.001187    0.007887  -0.151 0.880340
## Family3:Income     0.092702    0.021798   4.253 2.11e-05 ***
## Family4:Income     0.106746    0.021944   4.864 1.15e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2162.72  on 3494  degrees of freedom
## Residual deviance:  548.51  on 3477  degrees of freedom
## AIC: 584.51
##
## Number of Fisher Scoring iterations: 11
```

```
AIC(finalmodel2a)      # AIC = 584.5078
```

```
## [1] 584.5078
```

```
BIC(finalmodel2a)      # BIC = 695.3714
```

```
## [1] 695.3714
```

```
## Testing finalmodel2a
pred.final2a = predict(finalmodel2a, final_test, type = "response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
final2a.cutoff = 0.28
class.final2a = as.factor(if_else(pred.final2a < final2a.cutoff, "No", "Yes"))
prop.table(table(class.final2a))
```

```
## class.final2a
##      No      Yes
## 0.886 0.114
```

```
## Confusion Matrix
confusionMatrix(class.final2a, final.test$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##      No  1316  13
##      Yes   29  142
##
##           Accuracy : 0.972
##           95% CI : (0.9623, 0.9797)
##      No Information Rate : 0.8967
##      P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.8555
##
##      McNemar's Test P-Value : 0.02064
##
##           Sensitivity : 0.9784
##           Specificity : 0.9161
##      Pos Pred Value : 0.9902
##      Neg Pred Value : 0.8304
##           Prevalence : 0.8967
##      Detection Rate : 0.8773
##      Detection Prevalence : 0.8860
##      Balanced Accuracy : 0.9473
##
##      'Positive' Class : No
##
```

```
# Threshold = 0.28
# Accuracy = 0.972
# Sensitivity = 0.91613
# Specificity = 0.97844
```

```
## prefinalmodel2b: Repeated K-Fold Cross Validation with logged Mortgage, 3rd power Income, and Income*Family.
ctrl <- trainControl(method = "repeatedcv", number = 10, savePredictions = TRUE)
mod_fit <- train(Personal.Loan ~ poly(Income, 3) + Family + Income*Family + CCAvg + Education + Securities.Account
t + CD.Account + Online + CreditCard + LoggedMortgage,
  data = final_df,
  method = "glm", family = "binomial",
  trControl = ctrl, tuneLength = 5)
```

[illegible]

```
pred.final2b = predict(mod_fit, newdata = final_test)
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :  
## prediction from a rank-deficient fit may be misleading
```

```
confusionMatrix(data=pred.final2b, final_test$Personal.Loan)
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  No  Yes  
##           No 1336  23  
##           Yes   9 132  
##  
##           Accuracy : 0.9787  
##           95% CI : (0.97, 0.9854)  
##           No Information Rate : 0.8967  
##           P-Value [Acc > NIR] : < 2e-16  
##  
##           Kappa : 0.8801  
##  
##           McNemar's Test P-Value : 0.02156  
##  
##           Sensitivity : 0.9933  
##           Specificity : 0.8516  
##           Pos Pred Value : 0.9831  
##           Neg Pred Value : 0.9362  
##           Prevalence : 0.8967  
##           Detection Rate : 0.8907  
##           Detection Prevalence : 0.9060  
##           Balanced Accuracy : 0.9225  
##  
##           'Positive' Class : No  
##
```

```
# Accuracy    = 0.9787  
# Sensitivity = 0.8516  
# Specificity = 0.9933
```

```
## ROC Curve  
test_label = final_df[-split, "Personal.Loan"]  
results.modella = prediction(pred.finalla, test_label)  
length(pred.finalla)
```

```
## [1] 1500
```

```
length(test_label)
```

```
## [1] 1500
```

```
roc.modella = performance(results.modella, measure = "tpr", x.measure = "fpr")  
  
results.model2a = prediction(pred.final2a, test_label)  
length(pred.final2a)
```

```
## [1] 1500
```


```
length(test_label)
```

```
## [1] 1500
```



```
roc.model2a = performance(results.model2a, measure = "tpr", x.measure = "fpr")

plot(roc.step, col = "red", xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.poly.income3, col = "green", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.model1a, col = "blue", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.model2a, col = "orange", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
legend("bottomright", legend = c("Stepwise", "3rd Power Income", "Model 1A", "Model 2A"),
      col = c("red", "green", "blue", "orange"),
      lty=1, lwd=1)
```



Objective 2: LDA (continuous predictors only)

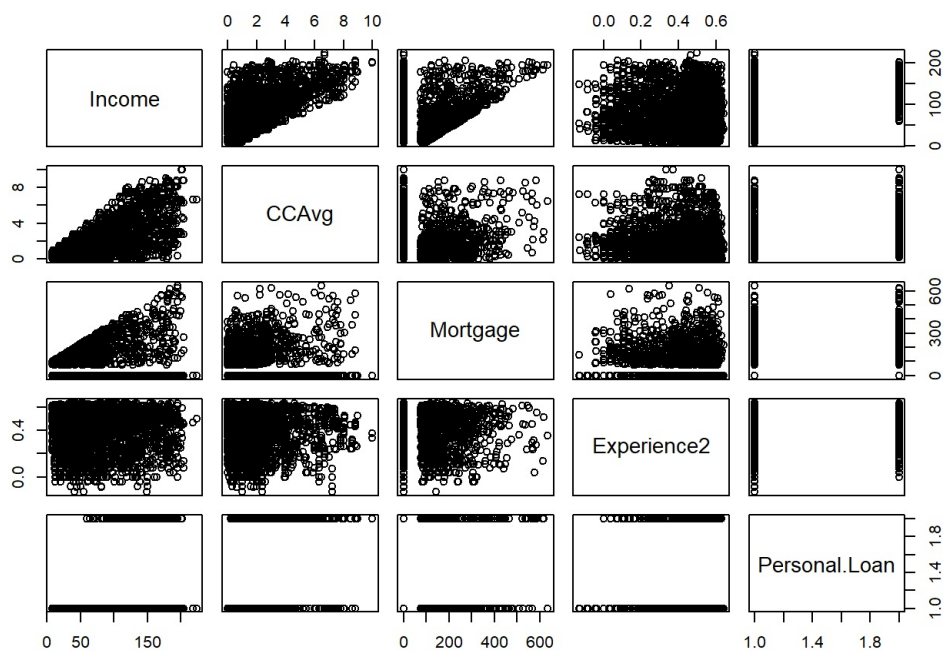
```
# Find only continuous predictors
str(train)
```

```
## 'data.frame': 3500 obs. of 11 variables:
## $ Income : int 23 52 98 79 95 63 91 143 59 38 ...
## $ Family : Factor w/ 4 levels "1","2","3","4": 3 4 1 2 2 4 1 3 3 1 ...
## $ CCAvg : num 0.4 1.3 5.4 2.8 0 3.6 0.1 2.9 0.9 1.5 ...
## $ Education : Factor w/ 3 levels "1","2","3": 1 2 1 1 3 3 2 3 3 2 ...
## $ Mortgage : int 0 0 0 179 0 0 199 0 199 116 ...
## $ Personal.Loan : Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 2 1 1 ...
## $ Securities.Account: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 2 1 ...
## $ CD.Account : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ Online : Factor w/ 2 levels "0","1": 2 2 2 1 2 1 2 2 2 1 ...
## $ CreditCard : Factor w/ 2 levels "0","1": 1 2 1 1 1 1 2 1 1 2 ...
## $ Experience2 : num 0.538 0.613 0.04 0.594 0.636 ...
```

```
# Setting up for PCA then LDA
LDA_train = select_if(train, is.numeric) %>% mutate(Personal.Loan = train$Personal.Loan)
str(LDA_train)
```

```
## 'data.frame': 3500 obs. of 5 variables:
## $ Income : int 23 52 98 79 95 63 91 143 59 38 ...
## $ CCAvg : num 0.4 1.3 5.4 2.8 0 3.6 0.1 2.9 0.9 1.5 ...
## $ Mortgage : int 0 0 0 179 0 0 199 0 199 116 ...
## $ Experience2 : num 0.538 0.613 0.04 0.594 0.636 ...
## $ Personal.Loan: Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 2 1 1 ...
```

```
pairs(LDA_train)
```




```
# PCA
reduced = LDA_train[-c(5)]
pc.result<-prcomp(reduced,scale=FALSE)
eigenvals<- (pc.result$sdev)^2
eigenvals
```


```
## [1] 1.038109e+04 1.988352e+03 1.730318e+00 2.731713e-02
```

```
plot(1:4,eigenvals/sum(eigenvals),type="l",main="Scree Plot",ylab="Prop. Var.
Explained")
```

Scree Plot



```
cumulative.prop<-cumsum(eigenvals/sum(eigenvals))
plot(1:4,cumulative.prop,type="l",main="Cumulative proportion",ylim=c(0,1))
```



```
par(mfrow=c(1,1))
```

```
# The desired number of PCs looks to be 1, since 2 retains 0% of the total variation.
```

```
# Build Model
```

```
LDA.model = lda(Personal.Loan ~ ., LDA_train)
```

```
LDA.model
```

```
## Call:
```

```
## lda(Personal.Loan ~ ., data = LDA_train)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##      No      Yes
```

```
## 0.90714286 0.09285714
```

```
##
```

```
## Group means:
```

```
##      Income      CCAvg      Mortgage      Experience2
```

```
## No   66.27969 1.738101  50.66677   0.4069236
```

```
## Yes 143.66154 3.803631 103.09846   0.3945928
```

```
##
```

```
## Coefficients of linear discriminants:
```

```
##      LD1
```

```
## Income      0.02262327
```

```
## CCAvg       0.07065000
```

```
## Mortgage    0.00119106
```

```
## Experience2 0.01861669
```

```
# Criteria
```

```
fit.p<-predict(LDA.model, newdata=test)
```

```
str(fit.p)
```

```
## List of 3
```

```
## $ class      : Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 1 1 1 ...
```

```
## $ posterior: num [1:1500, 1:2] 0.937 0.997 0.998 0.973 0.998 ...
```

```
## .. attr(*, "dimnames")=List of 2
```

```
## .. ..$ : chr [1:1500] "4" "6" "8" "9" ...
```

```
## .. ..$ : chr [1:2] "No" "Yes"
```

```
## $ x          : num [1:1500, 1] 0.586 -0.997 -1.344 0.132 -1.259 ...
```


```
## .. attr(*, "dimnames")=List of 2
```

```
## .. ..$ : chr [1:1500] "4" "6" "8" "9" ...
```

```
## .. ..$ : chr "LD1"
```

```
# ROC Curves
results.model<-prediction(fit.p$posterior[,2], test$Personal.Loan,label.ordering=c("No","Yes"))
roc.lda_ = performance(results.model, measure = "tpr", x.measure = "fpr")

plot(roc.step, col = "blue", xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.model2a, col = "red", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
plot(roc.lda_, col = "green", add = TRUE, xlim = c(0, 0.3), ylim = c(0.7, 1.0))
legend("bottomright", legend = c("Stepwise", "Complex", "LDA"),
      col = c("blue", "red", "green"),
      lty=1, lwd=1)
```



```
#fake<-train
#fake$Personal.Loan<-sample(fake$Personal.Loan,3500,replace=F)
#LDA.model.fake = lda(Personal.Loan ~ ., fake)
#LDA.model.fake

# Universally compare accuracy
confusionMatrix(class.step, test$Personal.Loan)$overall[1]
```

```
## Accuracy
## 0.968
```

```
confusionMatrix(class.final2a, test$Personal.Loan)$overall[1]
```

```
## Accuracy
## 0.972
```

```
round((mean(predict(LDA.model,newdata=test)$class==test$Personal.Loan)),3)
```

```
## [1] 0.899
```